

csm148finalproj

June 1, 2021

```
[1]: #Importing the necessary modules
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt # this is used for the plot the graph
import os
import seaborn as sns # used for plot interactive graph.
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn import metrics
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.cluster import KMeans
from sklearn.metrics import confusion_matrix
import sklearn.metrics.cluster as smc
from sklearn.model_selection import KFold

from matplotlib import pyplot
import itertools

%matplotlib inline

import random

random.seed(42)
```

```
[2]: def draw_confusion_matrix(y, yhat, classes):
    """
        Draws a confusion matrix for the given target and predictions
        Adapted from scikit-learn and discussion example.
    """
    plt.cla()
    plt.clf()
    matrix = confusion_matrix(y, yhat)
    plt.imshow(matrix, interpolation='nearest', cmap=plt.cm.Blues)
```

```

plt.title("Confusion Matrix")
plt.colorbar()
num_classes = len(classes)
plt.xticks(np.arange(num_classes), classes, rotation=90)
plt.yticks(np.arange(num_classes), classes)

fmt = 'd'
thresh = matrix.max() / 2.
for i, j in itertools.product(range(matrix.shape[0]), range(matrix.
→shape[1])):
    plt.text(j, i, format(matrix[i, j], fmt),
             horizontalalignment="center",
             color="white" if matrix[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
plt.show()

```

1 1. Loading the Dataset and Basic Stats

```
[3]: df_main = pd.read_csv('csm148finalprojectdata.csv')
```

```
[4]: df_main.head()
```

```
[4]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	51676	Female	61.0	0	0	Yes	
2	31112	Male	80.0	0	1	Yes	
3	60182	Female	49.0	0	0	Yes	
4	1665	Female	79.0	1	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Self-employed	Rural	202.21	NaN	never smoked	
2	Private	Rural	105.92	32.5	never smoked	
3	Private	Urban	171.23	34.4	smokes	
4	Self-employed	Rural	174.12	24.0	never smoked	

	stroke
0	1
1	1
2	1
3	1
4	1

```
[5]: #basic stats
df_main.describe()
```

```
[5]:
```

	id	age	hypertension	heart_disease \
count	5110.000000	5110.000000	5110.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012
std	21161.721625	22.612647	0.296607	0.226063
min	67.000000	0.080000	0.000000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000
max	72940.000000	82.000000	1.000000	1.000000

	avg_glucose_level	bmi	stroke
count	5110.000000	4909.000000	5110.000000
mean	106.147677	28.893237	0.048728
std	45.283560	7.854067	0.215320
min	55.120000	10.300000	0.000000
25%	77.245000	23.500000	0.000000
50%	91.885000	28.100000	0.000000
75%	114.090000	33.100000	0.000000
max	271.740000	97.600000	1.000000

```
[6]: df_main.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
id                5110 non-null int64
gender            5110 non-null object
age              5110 non-null float64
hypertension      5110 non-null int64
heart_disease     5110 non-null int64
ever_married      5110 non-null object
work_type         5110 non-null object
Residence_type    5110 non-null object
avg_glucose_level 5110 non-null float64
bmi               4909 non-null float64
smoking_status    5110 non-null object
stroke            5110 non-null int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.1+ KB
```

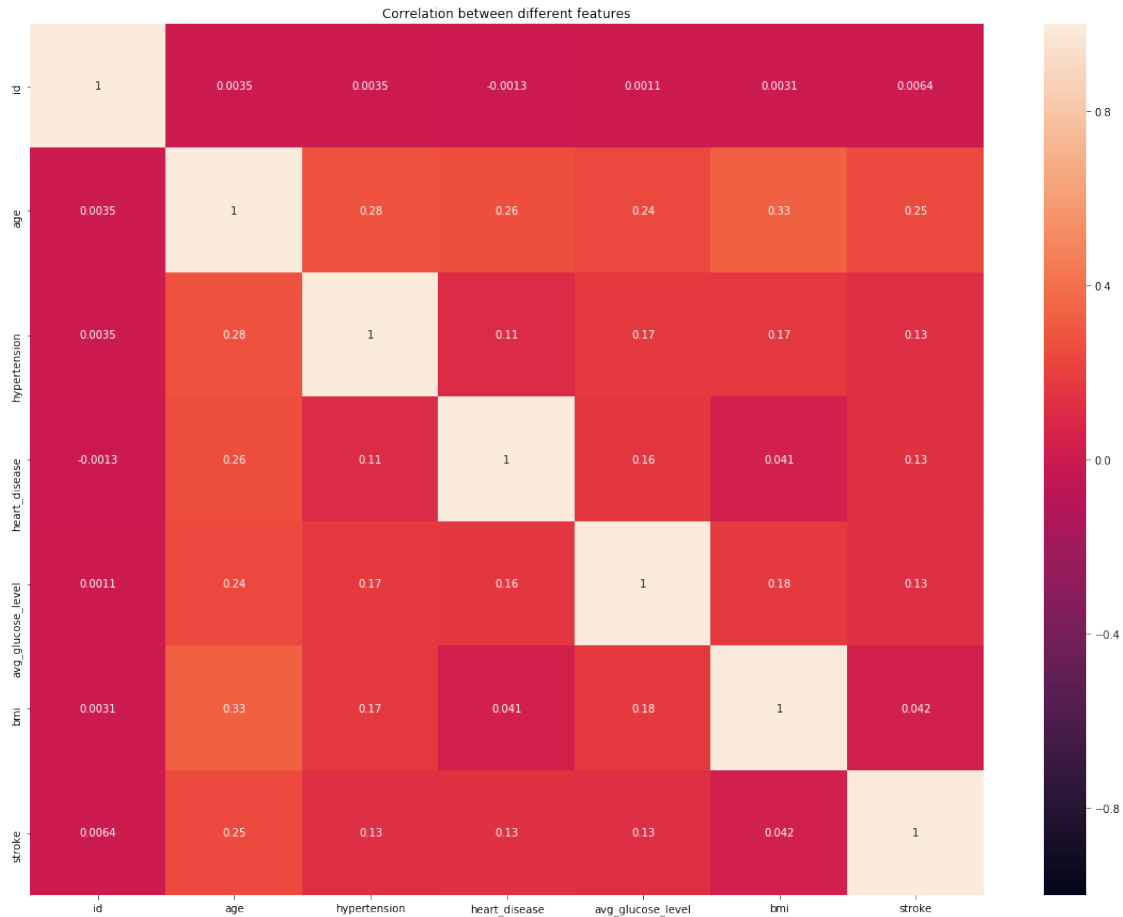
```
[7]: df_main.shape
```

```
[7]: (5110, 12)
```

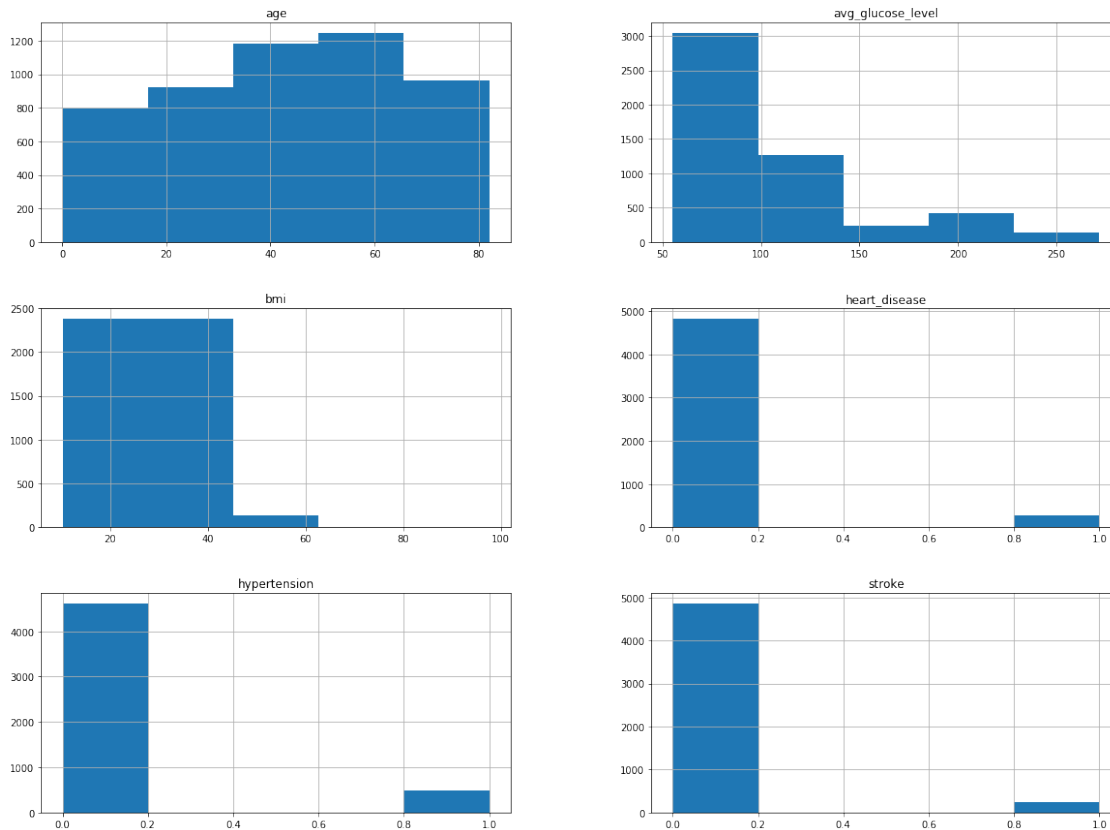
```
[8]: #Correlation between different labels
plt.figure(figsize=(20, 15))
```

```
hmap = sns.heatmap(df_main.corr(), vmin=-1, vmax=1, annot=True)
hmap.set_title('Correlation between different features')
```

[8]: Text(0.5, 1, 'Correlation between different features')



```
[9]: df_main = df_main.drop("id", axis=1)
df_main.hist(bins=5, figsize=(20,15))
plt.show()
```



2. Pipeline and Data Augmentation

```
[10]: from sklearn.compose import ColumnTransformer
      from sklearn.preprocessing import StandardScaler
      from sklearn.preprocessing import OneHotEncoder
      from sklearn.preprocessing import LabelEncoder
      from sklearn.pipeline import Pipeline
```

```
[11]: df1 = df_main.copy()
```

```
[12]: df1.columns
```

```
[12]: Index(['gender', 'age', 'hypertension', 'heart_disease', 'ever_married',
            'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
            'smoking_status', 'stroke'],
            dtype='object')
```

2.1 Determining categorization strategy

2.1.1 OHE - one hot encoding

2.1.2 LE - label encoding

```
[13]: df1['work_type'].unique() #OHE
[13]: array(['Private', 'Self-employed', 'Govt_job', 'children', 'Never_worked'],
          dtype=object)
[14]: df1['Residence_type'].unique() #LE
[14]: array(['Urban', 'Rural'], dtype=object)
[15]: df1['smoking_status'].unique() #OHE
[15]: array(['formerly smoked', 'never smoked', 'smokes', 'Unknown'],
          dtype=object)
[16]: df1['ever_married'].unique() #LE
[16]: array(['Yes', 'No'], dtype=object)
[17]: df1['gender'].unique() #OHE
[17]: array(['Male', 'Female', 'Other'], dtype=object)
```

2.2 Adding a new feature

```
[18]: #This was done initially to determine what feature to augment and commented out
      →on rerunning
      #it is done from scratch in the pipeline
      #df1['bmi_x_age'] = df1["bmi"]*df1["age"]
[19]: #corr_matrix = df1.corr()
      #corr_matrix["stroke"].sort_values(ascending=False)
[20]: #df1.describe()
```

3 Building the Pipeline

```
[21]: df1_features = df1.drop("stroke", axis=1)
      df1_labels = df1["stroke"].copy()
[22]: #Label encoding binary features prior to pipeline transformation
      le = LabelEncoder()
      for col in ["Residence_type", "ever_married"]:
          df1_features[col] = le.fit_transform(df1_features[col])
[23]: df1_features.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 10 columns):
gender                5110 non-null object
age                  5110 non-null float64
hypertension         5110 non-null int64
heart_disease        5110 non-null int64
ever_married         5110 non-null int64
work_type            5110 non-null object
Residence_type       5110 non-null int64
avg_glucose_level    5110 non-null float64
bmi                  4909 non-null float64
smoking_status       5110 non-null object
dtypes: float64(3), int64(4), object(3)
memory usage: 399.3+ KB

```

```
[24]: df1_features.head()
```

```

[24]:   gender  age  hypertension  heart_disease  ever_married  work_type \
0   Male  67.0             0             1             1   Private
1  Female  61.0             0             0             1  Self-employed
2   Male  80.0             0             1             1   Private
3  Female  49.0             0             0             1   Private
4  Female  79.0             1             0             1  Self-employed

      Residence_type  avg_glucose_level  bmi  smoking_status
0                1          228.69  36.6  formerly smoked
1                0          202.21   NaN    never smoked
2                0          105.92  32.5    never smoked
3                1          171.23  34.4         smokes
4                0          174.12  24.0    never smoked

```

```

[25]: # This cell implements the complete pipeline for preparing the data
from sklearn.impute import SimpleImputer
from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="mean") # use mean imputation for missing
    → values
# remove the categorical features
df1_num = df1_features.drop(["gender", "ever_married", "heart_disease",
    → "hypertension", "work_type", "Residence_type", "smoking_status"], axis=1)

age_idx, bmi_idx = 0 , 2

class AugmentFeatures(BaseEstimator, TransformerMixin):
    '''
    implements the previous features we had defined

```

```

df1['bmi_x_age'] = df1["bmi"]*df1["age"]
'''
def __init__(self, add_bmi_x_age = True):
    self.add_bmi_x_age = add_bmi_x_age

def fit(self, X, y=None):
    return self # nothing else to do

def transform(self, X):
    final = X
    if self.add_bmi_x_age:
        bmi_x_age = X[:, bmi_idx] * X[:, age_idx]
        final = np.c_[X, bmi_x_age]
    return final

attr_adder = AugmentFeatures()
df1_extra_attribs = attr_adder.transform(df1.values) # generate new features

# this will be a numerical pipeline
# 1. impute, 2. augment the feature set 3. normalize using StandardScaler()
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="mean")),
    ('attribs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

df1_num_tr = num_pipeline.fit_transform(df1_num)

numerical_features = list(df1_num)
categorical_features = ["gender", "work_type", "smoking_status"]

full_pipeline = ColumnTransformer(transformers=[
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features)],
    remainder='passthrough')

data_prepared = full_pipeline.fit_transform(df1_features)

```

[26]: data_prepared

```

[26]: array([[ 1.05143428e+00,  2.70637544e+00,  1.00123401e+00, ...,
                1.00000000e+00,  1.00000000e+00,  1.00000000e+00],
               [ 7.86070073e-01,  2.12155854e+00,  4.6155355e-16, ...,
                0.00000000e+00,  1.00000000e+00,  0.00000000e+00],
               [ 1.62639008e+00, -5.02830130e-03,  4.68577254e-01, ...,
                1.00000000e+00,  1.00000000e+00,  0.00000000e+00],
               ...,
               [-3.63841511e-01, -5.11442636e-01,  2.21736316e-01, ...,

```



```

0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[ 3.43796387e-01, 1.32825706e+00, -4.27845098e-01, ...,
0.00000000e+00, 1.00000000e+00, 0.00000000e+00],
[ 3.42048064e-02, -4.60867458e-01, -3.49895329e-01, ...,
0.00000000e+00, 1.00000000e+00, 1.00000000e+00]])

```

3.1 Splitting the Dataset

```

[27]: Xp_train, Xp_test, Yp_train, Yp_test = train_test_split(data_prepared,
    ↪ df1_labels, test_size=0.2, random_state=42)
print("Xp_train shape:", Xp_train.shape)
print("Yp_train shape:", Yp_train.shape)
print("Xp_test shape:", Xp_test.shape)
print("Yp_test shape:", Yp_test.shape)

```

```

Xp_train shape: (4088, 20)
Yp_train shape: (4088,)
Xp_test shape: (1022, 20)
Yp_test shape: (1022,)

```

```

[28]: Yp_train.value_counts()

```

```

[28]: 0    3901
      1    187
      Name: stroke, dtype: int64

```

3.2 Balancing the dataset

```

[29]: from imblearn.over_sampling import SMOTE

```

```

[30]: sm = SMOTE(random_state=20)
      Xp_train, Yp_train = sm.fit_resample(Xp_train, Yp_train)

```

```

[31]: Yp_train.value_counts()

```

```

[31]: 1    3901
      0    3901
      Name: stroke, dtype: int64

```

4 3. Logistic Regression

```

[32]: from sklearn.metrics import precision_score
      from sklearn.metrics import recall_score
      from sklearn.metrics import f1_score
      from sklearn.metrics import accuracy_score

```

4.0.1 solver='lbfgs'

```
[33]: log_clf = LogisticRegression(solver='lbfgs').fit(Xp_train, Yp_train)
      y_pred = log_clf.predict(Xp_test)
```

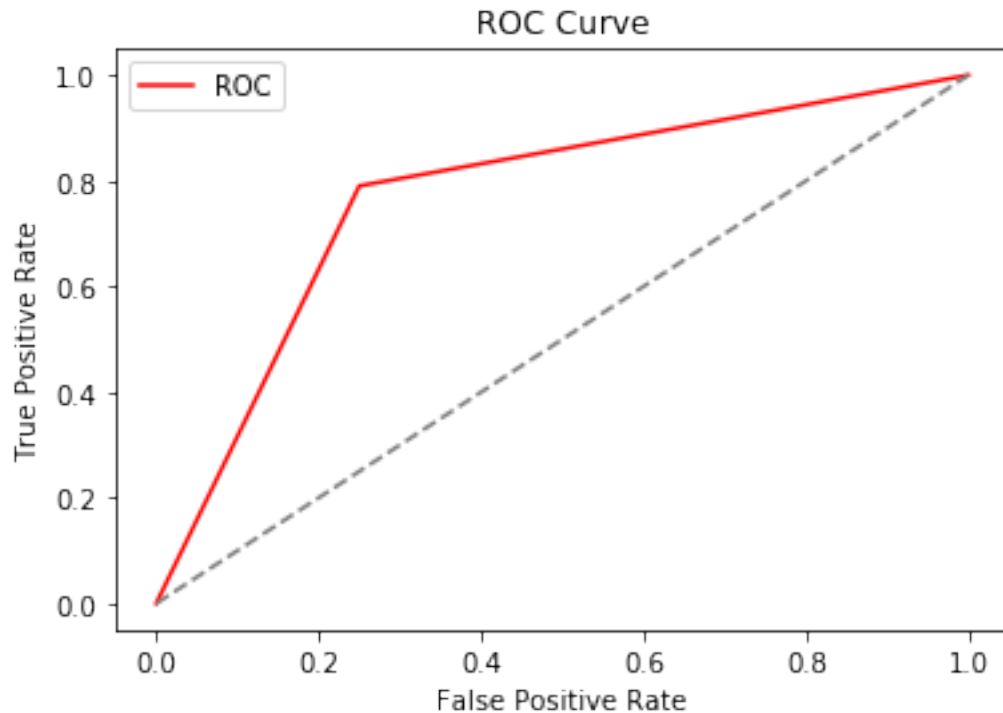
```
[34]: acc_log = accuracy_score(Yp_test, y_pred)
      prec_log = precision_score(Yp_test, y_pred)
      rec_log = recall_score(Yp_test, y_pred)
      f1_log = f1_score(Yp_test, y_pred)

      print("Accuracy:", acc_log)
      print("Precision:", prec_log)
      print("Recall:", rec_log)
      print("F1 Score:", f1_log)
```

Accuracy: 0.7524461839530333
Precision: 0.1695501730103806
Recall: 0.7903225806451613
F1 Score: 0.2792022792022792

```
[35]: fpr, tpr, threshold = metrics.roc_curve(Yp_test, y_pred)
      plt.plot(fpr, tpr, color='red', label='ROC')
      plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.legend()
      plt.show()

      auc = np.trapz(tpr, fpr)
      print('AUC:', auc)
```



AUC: 0.7701612903225806

4.0.2 solver='sag'

```
[36]: log_clf2 = LogisticRegression(penalty='none', solver='sag', max_iter=100).
      → fit(Xp_train, Yp_train)
      y_pred1 = log_clf2.predict(Xp_test)
```

```
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/linear_model/_sag.py:329: ConvergenceWarning: The max_iter was
reached which means the coef_ did not converge
  "the coef_ did not converge", ConvergenceWarning)
```

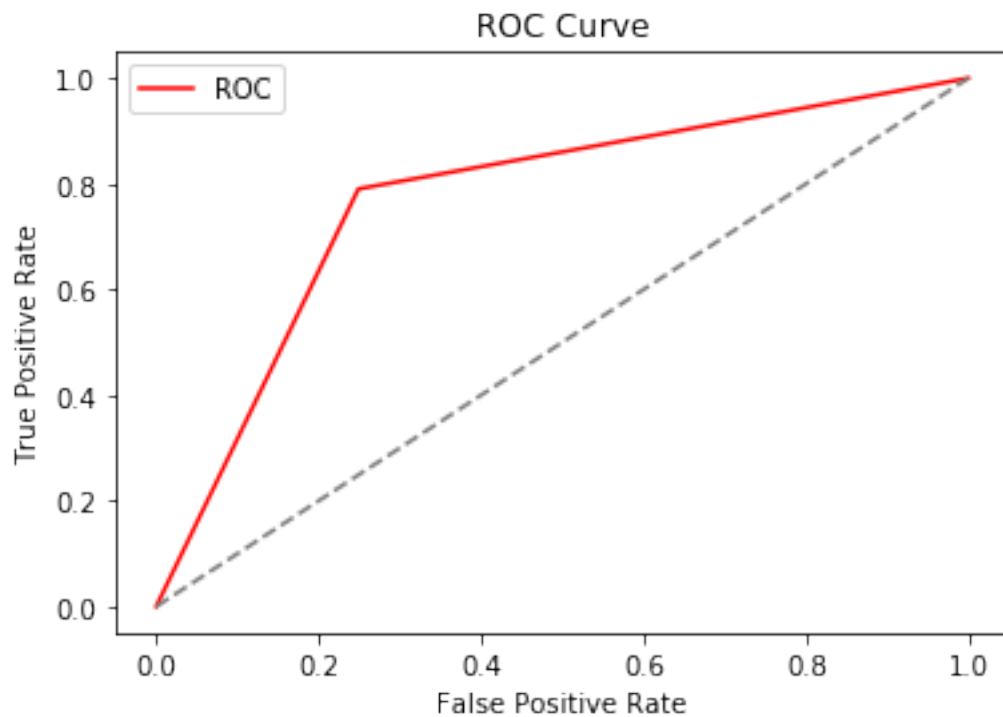
```
[37]: acc_log = accuracy_score(Yp_test, y_pred1)
      prec_log = precision_score(Yp_test, y_pred1)
      rec_log = recall_score(Yp_test, y_pred1)
      f1_log = f1_score(Yp_test, y_pred1)

      print("Accuracy:", acc_log)
      print("Precision:", prec_log)
      print("Recall:", rec_log)
      print("F1 Score:", f1_log)
```

Accuracy: 0.7534246575342466
Precision: 0.1701388888888889
Recall: 0.7903225806451613
F1 Score: 0.27999999999999997

```
[38]: fpr, tpr, threshold = metrics.roc_curve(Yp_test, y_pred1)
plt.plot(fpr, tpr, color='red', label='ROC')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

auc = np.trapz(tpr,fpr)
print('AUC:', auc)
```



AUC: 0.7706821236559139

4.0.3 Associated p-values and regression

```
[39]: import statsmodels.api as sms
```

```
[40]: # build the OLS model (ordinary least squares) from the training data
stroke_stats = sms.OLS(df1_labels, data_prepared)

# do the fit and save regression info (parameters, etc) in results_stats
results_stats = stroke_stats.fit()

[41]: print(results_stats.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  stroke    R-squared:                  0.085
Model:                          OLS      Adj. R-squared:             0.082
Method:                        Least Squares    F-statistic:                27.73
Date:                          Tue, 01 Jun 2021    Prob (F-statistic):        3.07e-85
Time:                          12:34:47      Log-Likelihood:            822.99
No. Observations:              5110      AIC:                      -1610.
Df Residuals:                  5092      BIC:                      -1492.
Df Model:                      17
Covariance Type:               nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
x1	0.0857	0.015	5.696	0.000	0.056	0.115
x2	0.0141	0.003	4.593	0.000	0.008	0.020
x3	0.0030	0.008	0.379	0.705	-0.012	0.018
x4	-0.0187	0.017	-1.107	0.268	-0.052	0.014
x5	0.0352	0.040	0.880	0.379	-0.043	0.114
x6	0.0337	0.040	0.840	0.401	-0.045	0.112
x7	0.0092	0.167	0.055	0.956	-0.318	0.337
x8	-0.0078	0.020	-0.384	0.701	-0.048	0.032
x9	0.0305	0.042	0.731	0.465	-0.051	0.112
x10	0.0067	0.019	0.347	0.728	-0.031	0.044
x11	-0.0127	0.020	-0.626	0.532	-0.052	0.027
x12	0.0614	0.022	2.796	0.005	0.018	0.104
x13	0.0215	0.023	0.939	0.348	-0.023	0.066
x14	0.0227	0.023	0.992	0.321	-0.022	0.068
x15	0.0138	0.023	0.609	0.542	-0.031	0.058
x16	0.0200	0.023	0.861	0.389	-0.026	0.066
x17	0.0389	0.010	3.776	0.000	0.019	0.059
x18	0.0501	0.013	3.714	0.000	0.024	0.077
x19	-0.0348	0.009	-4.061	0.000	-0.052	-0.018
x20	0.0052	0.006	0.907	0.365	-0.006	0.017

```

=====
Omnibus:                      3802.843    Durbin-Watson:              0.172
Prob(Omnibus):                 0.000      Jarque-Bera (JB):           47486.467
Skew:                          3.646      Prob(JB):                   0.00
Kurtosis:                     16.032      Cond. No.                   1.88e+16
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.92e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

5 4. PCA

```
[42]: from sklearn import decomposition
```

```
[43]: pca_data = data_prepared.copy()
```

```
[44]: pca_data
```

```
[44]: array([[ 1.05143428e+00,  2.70637544e+00,  1.00123401e+00, ...,
           1.00000000e+00,  1.00000000e+00,  1.00000000e+00],
        [ 7.86070073e-01,  2.12155854e+00,  4.61555355e-16, ...,
           0.00000000e+00,  1.00000000e+00,  0.00000000e+00],
        [ 1.62639008e+00, -5.02830130e-03,  4.68577254e-01, ...,
           1.00000000e+00,  1.00000000e+00,  0.00000000e+00],
        ...,
        [-3.63841511e-01, -5.11442636e-01,  2.21736316e-01, ...,
           0.00000000e+00,  1.00000000e+00,  0.00000000e+00],
        [ 3.43796387e-01,  1.32825706e+00, -4.27845098e-01, ...,
           0.00000000e+00,  1.00000000e+00,  0.00000000e+00],
        [ 3.42048064e-02, -4.60867458e-01, -3.49895329e-01, ...,
           0.00000000e+00,  1.00000000e+00,  1.00000000e+00]])
```

```
[45]: pca = decomposition.PCA(0.8)

      # Now we run the fit operation to convert our
      # data to a PCA transformed data
pca_data = pca.fit_transform(pca_data)
```

```
[46]: pca_data.shape
```

```
[46]: (5110, 6)
```

5.1 Splitting the data

```
[47]: new_X_train, new_X_test, new_Y_train, new_Y_test = train_test_split(pca_data,
    ↪ df1_labels, test_size=0.2, random_state=42)
```

5.2 Balancing the Data

```
[48]: sm = SMOTE(random_state=20)
      new_X_train, new_Y_train = sm.fit_resample(new_X_train, new_Y_train)
```

```
[49]: new_Y_train.value_counts()
```

```
[49]: 1    3901
      0    3901
      Name: stroke, dtype: int64
```

5.3 Implementing Logistic Regression after PCA transformation

5.4 solver='lbfgs'

```
[50]: log_clf = LogisticRegression(solver='lbfgs').fit(new_X_train, new_Y_train)
      y_pred = log_clf.predict(new_X_test)
```

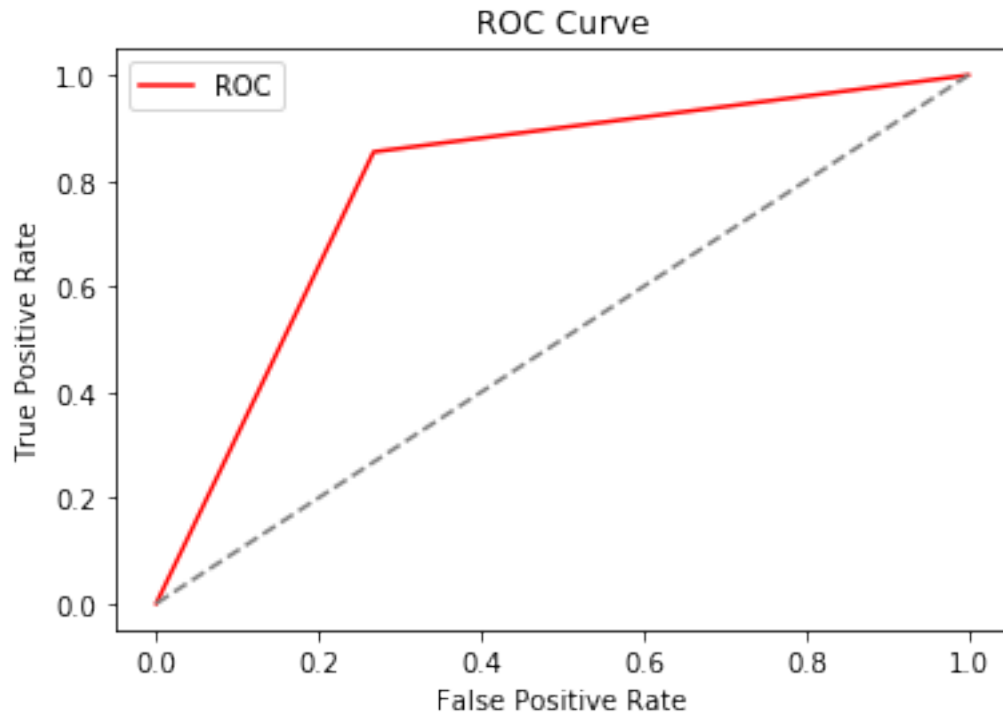
```
[51]: acc_log = accuracy_score(new_Y_test, y_pred)
      prec_log = precision_score(new_Y_test, y_pred)
      rec_log = recall_score(new_Y_test, y_pred)
      f1_log = f1_score(new_Y_test, y_pred)
```

```
print("Accuracy:", acc_log)
print("Precision:", prec_log)
print("Recall:", rec_log)
print("F1 Score:", f1_log)
```

```
Accuracy: 0.7397260273972602
Precision: 0.17096774193548386
Recall: 0.8548387096774194
F1 Score: 0.2849462365591398
```

```
[52]: fpr, tpr, threshold = metrics.roc_curve(Yp_test, y_pred)
      plt.plot(fpr, tpr, color='red', label='ROC')
      plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('ROC Curve')
      plt.legend()
      plt.show()

      auc = np.trapz(tpr, fpr)
      print('AUC:', auc)
```



AUC: 0.793565188172043

5.5 solver='sag'

```
[53]: log_clf2 = LogisticRegression(penalty='none', solver='sag', max_iter=100).
      → fit(new_X_train, new_Y_train)
      y_pred1 = log_clf2.predict(new_X_test)
```

```
[54]: acc_log = accuracy_score(new_Y_test, y_pred1)
      prec_log = precision_score(new_Y_test, y_pred1)
      rec_log = recall_score(new_Y_test, y_pred1)
      f1_log = f1_score(new_Y_test, y_pred1)

      print("Accuracy:", acc_log)
      print("Precision:", prec_log)
      print("Recall:", rec_log)
      print("F1 Score:", f1_log)
```

Accuracy: 0.7397260273972602

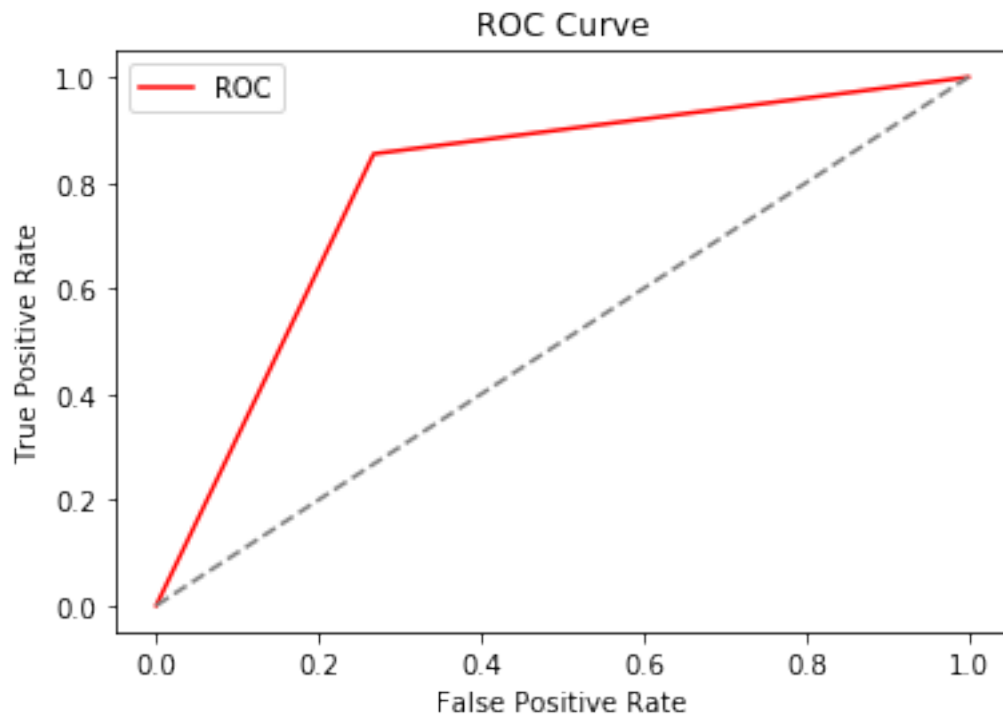
Precision: 0.17096774193548386

Recall: 0.8548387096774194

F1 Score: 0.2849462365591398


```
[55]: fpr, tpr, threshold = metrics.roc_curve(Yp_test, y_pred1)
plt.plot(fpr, tpr, color='red', label='ROC')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

auc = np.trapz(tpr,fpr)
print('AUC:', auc)
```



AUC: 0.793565188172043

6 5. Bagging

```
[56]: from sklearn.ensemble import BaggingClassifier
from sklearn import model_selection
```

```
[57]: tree = DecisionTreeClassifier(max_depth=3, random_state=20)
```

6.1 Using non-PCA data

```
[58]: bagging = BaggingClassifier(base_estimator=tree, n_estimators=16,
    ↳max_samples=200, bootstrap=True)

[59]: bagging.fit(Xp_train, Yp_train)

[59]: BaggingClassifier(base_estimator=DecisionTreeClassifier(max_depth=3,
    random_state=20),
    max_samples=200, n_estimators=16)

[60]: y_pred = bagging.predict(Xp_test)

[61]: acc_bag = accuracy_score(Yp_test, y_pred)
    prec_bag = precision_score(Yp_test, y_pred)
    rec_bag = recall_score(Yp_test, y_pred)
    f1_bag = f1_score(Yp_test, y_pred)

    print("Accuracy:", acc_bag)
    print("Precision:", prec_bag)
    print("Recall:", rec_bag)
    print("F1 Score:", f1_bag)
```

```
Accuracy: 0.7397260273972602
Precision: 0.16883116883116883
Recall: 0.8387096774193549
F1 Score: 0.28108108108108104
```

6.2 Using PCA data

```
[62]: bagging2 = BaggingClassifier(base_estimator=tree, n_estimators=6,
    ↳max_samples=200, bootstrap=True)

[63]: bagging2.fit(new_X_train, new_Y_train)

[63]: BaggingClassifier(base_estimator=DecisionTreeClassifier(max_depth=3,
    random_state=20),
    max_samples=200, n_estimators=6)

[64]: y_pred = bagging2.predict(new_X_test)

[65]: acc_bag = accuracy_score(new_Y_test, y_pred)
    prec_bag = precision_score(new_Y_test, y_pred)
    rec_bag = recall_score(new_Y_test, y_pred)
    f1_bag = f1_score(new_Y_test, y_pred)

    print("Accuracy:", acc_bag)
    print("Precision:", prec_bag)
    print("Recall:", rec_bag)
    print("F1 Score:", f1_bag)
```

Accuracy: 0.700587084148728
Precision: 0.1534090909090909
Recall: 0.8709677419354839
F1 Score: 0.2608695652173913

6.3 Hyperparameter tuning (using only PCA transformed data)

```
[77]: from sklearn.model_selection import GridSearchCV
      from sklearn.model_selection import RepeatedStratifiedKFold
      from sklearn.model_selection import cross_val_score
      from sklearn.metrics import f1_score, make_scorer

[78]: f1 = make_scorer(f1_score , average='macro')

[79]: n_estimators = [10, 50, 100]
      max_samples = [0.6, 0.8, 1.0]
      max_features = [4, 5, 6]

[80]: grid = dict(n_estimators = n_estimators, max_samples = max_samples,
                  max_features = max_features)
      grid_search = GridSearchCV(BaggingClassifier(base_estimator=DecisionTreeClassifier()),
      param_grid=grid, n_jobs=-1, cv=5, scoring=f1)

[81]: grid_result = grid_search.fit(new_X_train, new_Y_train)

[82]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

Best: 0.936358 using {'max_features': 4, 'max_samples': 1.0, 'n_estimators': 100}

7 6. Neural Network

Note: From this point on, only PCA tranformed data is used.

```
[72]: from sklearn.neural_network import MLPClassifier
      from sklearn.datasets import make_classification

[73]: #We only use PCA transformed data from this point on
      mlp_clf = MLPClassifier(hidden_layer_sizes=(150,100,50), max_iter=200, alpha=0.
      0001,
      solver='adam', verbose=10, random_state=21)
      mlp_clf.fit(new_X_train, new_Y_train)

      y_pred = mlp_clf.predict(new_X_test)
```

Iteration 1, loss = 0.51515358
Iteration 2, loss = 0.43941632
Iteration 3, loss = 0.42570626
Iteration 4, loss = 0.43331772

Iteration 5, loss = 0.40764382
Iteration 6, loss = 0.39441906
Iteration 7, loss = 0.38384440
Iteration 8, loss = 0.37244122
Iteration 9, loss = 0.36142105
Iteration 10, loss = 0.35324388
Iteration 11, loss = 0.33810572
Iteration 12, loss = 0.33745566
Iteration 13, loss = 0.31911395
Iteration 14, loss = 0.31699107
Iteration 15, loss = 0.32137112
Iteration 16, loss = 0.32485903
Iteration 17, loss = 0.31909150
Iteration 18, loss = 0.31202846
Iteration 19, loss = 0.28559606
Iteration 20, loss = 0.36002740
Iteration 21, loss = 0.31929477
Iteration 22, loss = 0.28196541
Iteration 23, loss = 0.26923592
Iteration 24, loss = 0.26142548
Iteration 25, loss = 0.28480182
Iteration 26, loss = 0.26371197
Iteration 27, loss = 0.27743039
Iteration 28, loss = 0.25938998
Iteration 29, loss = 0.24255022
Iteration 30, loss = 0.23457857
Iteration 31, loss = 0.23201297
Iteration 32, loss = 0.22786864
Iteration 33, loss = 0.21947472
Iteration 34, loss = 0.21545828
Iteration 35, loss = 0.21610703
Iteration 36, loss = 0.20792997
Iteration 37, loss = 0.21144203
Iteration 38, loss = 0.23039354
Iteration 39, loss = 0.22156244
Iteration 40, loss = 0.22882275
Iteration 41, loss = 0.33956993
Iteration 42, loss = 0.22452186
Iteration 43, loss = 0.21166971
Iteration 44, loss = 0.20159946
Iteration 45, loss = 0.21418709
Iteration 46, loss = 0.19603618
Iteration 47, loss = 0.19648629
Iteration 48, loss = 0.21354309
Iteration 49, loss = 0.18690936
Iteration 50, loss = 0.18104178
Iteration 51, loss = 0.18716763
Iteration 52, loss = 0.25230443

Iteration 53, loss = 0.18803876
Iteration 54, loss = 0.18229915
Iteration 55, loss = 0.22434972
Iteration 56, loss = 0.17428616
Iteration 57, loss = 0.17940882
Iteration 58, loss = 0.16849219
Iteration 59, loss = 0.16520146
Iteration 60, loss = 0.16377354
Iteration 61, loss = 0.17910413
Iteration 62, loss = 0.16859064
Iteration 63, loss = 0.16063379
Iteration 64, loss = 0.20140106
Iteration 65, loss = 0.22355943
Iteration 66, loss = 0.17052399
Iteration 67, loss = 0.16367491
Iteration 68, loss = 0.15560659
Iteration 69, loss = 0.33035073
Iteration 70, loss = 0.18592161
Iteration 71, loss = 0.16743508
Iteration 72, loss = 0.15996382
Iteration 73, loss = 0.21890874
Iteration 74, loss = 0.16932110
Iteration 75, loss = 0.15612446
Iteration 76, loss = 0.16724629
Iteration 77, loss = 0.15575383
Iteration 78, loss = 0.15087445
Iteration 79, loss = 0.30816145
Iteration 80, loss = 0.22119059
Iteration 81, loss = 0.24719840
Iteration 82, loss = 0.17406028
Iteration 83, loss = 0.16809002
Iteration 84, loss = 0.16255774
Iteration 85, loss = 0.20065668
Iteration 86, loss = 0.15877248
Iteration 87, loss = 0.15177169
Iteration 88, loss = 0.15704586
Iteration 89, loss = 0.14861864
Iteration 90, loss = 0.14333641
Iteration 91, loss = 0.13963085
Iteration 92, loss = 0.14870546
Iteration 93, loss = 0.13783409
Iteration 94, loss = 0.13758879
Iteration 95, loss = 0.30640821
Iteration 96, loss = 0.17582867
Iteration 97, loss = 0.21292823
Iteration 98, loss = 0.17478652
Iteration 99, loss = 0.15695555
Iteration 100, loss = 0.14857294

```

Iteration 101, loss = 0.14097499
Iteration 102, loss = 0.25251360
Iteration 103, loss = 0.15266756
Iteration 104, loss = 0.16144775
Iteration 105, loss = 0.14711195
Training loss did not improve more than tol=0.000100 for 10 consecutive epochs.
Stopping.

```

```

[74]: acc_nn = accuracy_score(new_Y_test, y_pred)
      prec_nn = precision_score(new_Y_test, y_pred)
      rec_nn = recall_score(new_Y_test, y_pred)
      f1_nn = f1_score(new_Y_test, y_pred)

      print("Accuracy:", acc_nn)
      print("Precision:", prec_nn)
      print("Recall:", rec_nn)
      print("F1 Score:", f1_nn)

```

```

Accuracy: 0.8424657534246576
Precision: 0.14388489208633093
Recall: 0.3225806451612903
F1 Score: 0.19900497512437812

```

7.1 Hyperparameter tuning

```

[84]: nn_clf = MLPClassifier(max_iter=200)

      parameter_space = {
          'hidden_layer_sizes': [(50,100,150), (100,)],
          'activation': ['tanh', 'relu'],
          'solver': ['sgd', 'adam'],
          'alpha': [0.0001, 0.001],
          'learning_rate': ['constant', 'adaptive'],
      }

      grid_clf = GridSearchCV(nn_clf, parameter_space, n_jobs=-1, cv=3, scoring=f1)

[85]: grid_result = grid_clf.fit(new_X_train, new_Y_train)

[86]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))

```

```

Best: 0.924051 using {'activation': 'relu', 'alpha': 0.0001,
'hidden_layer_sizes': (50, 100, 150), 'learning_rate': 'adaptive', 'solver':
'adam'}

```

8 7. K-Fold Cross Validation

```
[94]: from sklearn.model_selection import KFold
      from sklearn import model_selection
```

8.1 Stratified K Fold

```
[100]: kfold = model_selection.StratifiedKFold(n_splits=5)

[102]: bag_model_kfold = BaggingClassifier(base_estimator=tree, n_estimators=6,
      ↪max_samples=200, bootstrap=True)

nn_model_kfold = MLPClassifier(hidden_layer_sizes=(150,100,50), max_iter=200,
      ↪alpha=0.001,
      solver='adam')

bag_results_kfold = model_selection.cross_val_score(bag_model_kfold, pca_data,
      ↪df1_labels, cv=kfold, scoring=f1)

nn_results_kfold = model_selection.cross_val_score(nn_model_kfold, pca_data,
      ↪df1_labels, cv=kfold, scoring=f1)

# Because we're collecting results from all runs, we take the mean value
print(" Bagging f1 score: %.2f%%" % (bag_results_kfold.mean()*100.0))

print("Neural Network f1 score: %.2f%%" % (nn_results_kfold.mean()*100.0))
```

```
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
  % self.max_iter, ConvergenceWarning)

Bagging f1 score: 50.55%
Neural Network f1 score: 53.51%
```

```

/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)

```

8.2 StratifiedShuffleSplit

```

[103]: kfold = model_selection.StratifiedShuffleSplit(n_splits=5, test_size=0.2,
    ↪random_state=20)

[104]: bag_model_kfold = BaggingClassifier(base_estimator=tree, n_estimators=6,
    ↪max_samples=200, bootstrap=True)

nn_model_kfold = MLPClassifier(hidden_layer_sizes=(150,100,50), max_iter=200,
    ↪alpha=0.001,
                                solver='adam')

bag_results_kfold = model_selection.cross_val_score(bag_model_kfold, pca_data,
    ↪df1_labels, cv=kfold, scoring=f1)

nn_results_kfold = model_selection.cross_val_score(nn_model_kfold, pca_data,
    ↪df1_labels, cv=kfold, scoring=f1)

# Because we're collecting results from all runs, we take the mean value
print(" Bagging f1 score: %.2f%%" % (bag_results_kfold.mean()*100.0))

print("Neural Network f1 score: %.2f%%" % (nn_results_kfold.mean()*100.0))

```

```

/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)
/Users/ojasbardiya/anaconda3/lib/python3.7/site-
packages/sklearn/neural_network/_multilayer_perceptron.py:617:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and
the optimization hasn't converged yet.
    % self.max_iter, ConvergenceWarning)

```

```

Bagging f1 score: 49.84%
Neural Network f1 score: 52.10%

```



```
/Users/ojasbardiya/anaconda3/lib/python3.7/site-  
packages/sklearn/neural_network/_multilayer_perceptron.py:617:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and  
the optimization hasn't converged yet.  
    % self.max_iter, ConvergenceWarning)
```

8.3 8. Custom Models

8.4 SVM

```
[107]: clf_svm = SVC(probability=True, gamma='scale')  
clf_svm.fit(new_X_train, new_Y_train)
```

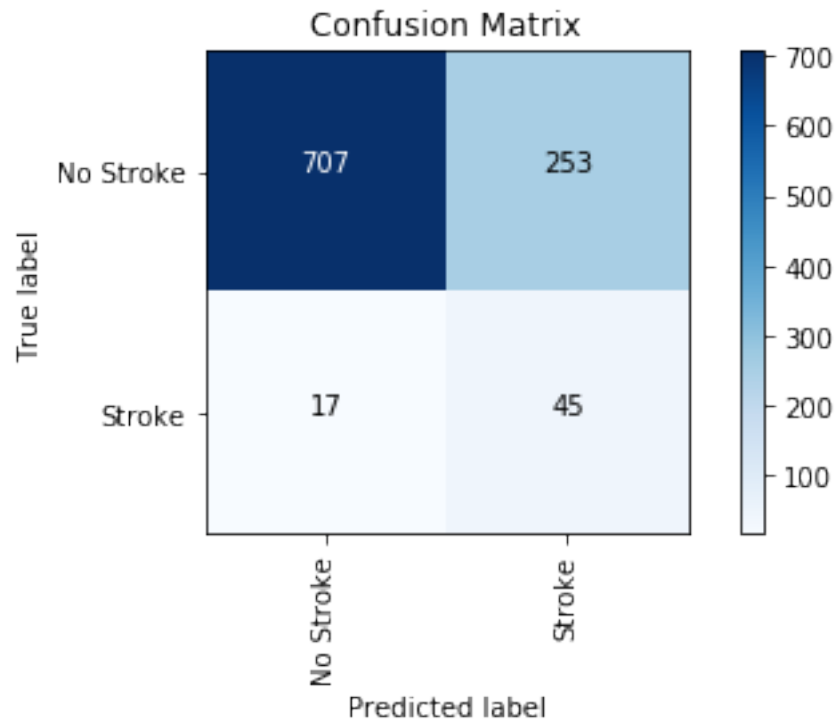
```
[107]: SVC(probability=True)
```

```
[109]: y_svm_pred = clf_svm.predict(new_X_test)
```

```
[113]: acc_svm = accuracy_score(new_Y_test, y_svm_pred)  
prec_svm = precision_score(new_Y_test, y_svm_pred)  
rec_svm = recall_score(new_Y_test, y_svm_pred)  
f1_svm = f1_score(new_Y_test, y_svm_pred)  
  
print("Accuracy:", acc_svm)  
print("Precision:", prec_svm)  
print("Recall:", rec_svm)  
print("F1 Score:", f1_svm)
```

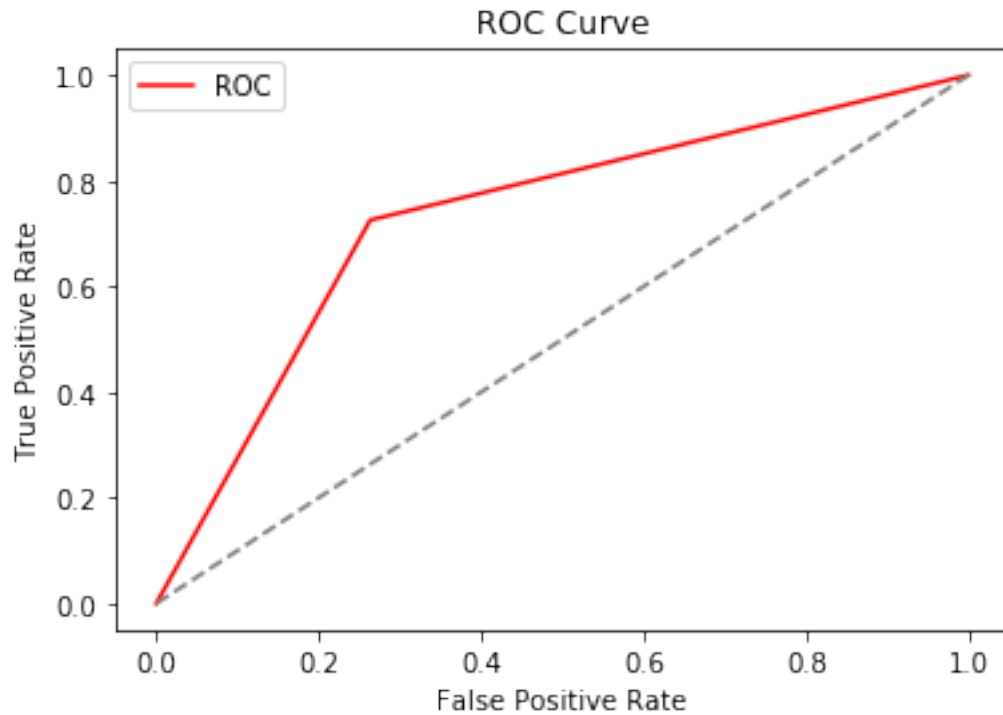
```
Accuracy: 0.735812133072407  
Precision: 0.15100671140939598  
Recall: 0.7258064516129032  
F1 Score: 0.25
```

```
[114]: draw_confusion_matrix(new_Y_test, y_svm_pred, ['No Stroke', 'Stroke'])
```



```
[115]: fpr, tpr, threshold = metrics.roc_curve(new_Y_test, y_svm_pred)
plt.plot(fpr, tpr, color='red', label='ROC')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()

auc = np.trapz(tpr,fpr)
print('AUC:', auc)
```



AUC: 0.7311323924731183

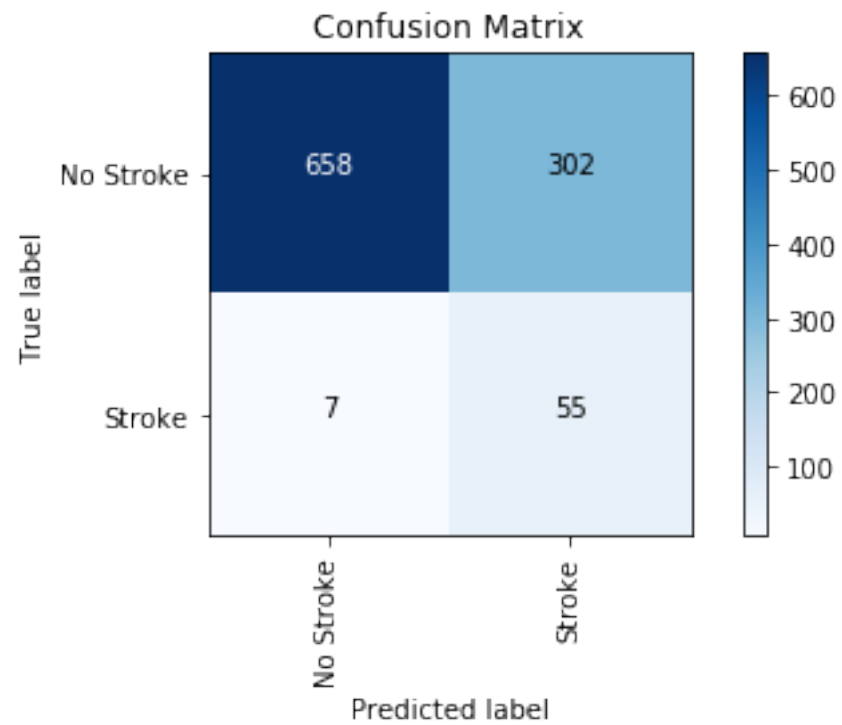
8.5 Bayesian Classification

```
[116]: from sklearn.naive_bayes import GaussianNB
[117]: clf_bayes = GaussianNB().fit(new_X_train, new_Y_train)
[118]: preds = clf_bayes.predict(new_X_test)
[119]: acc_bayes = accuracy_score(new_Y_test, preds)
      prec_bayes = precision_score(new_Y_test, preds)
      rec_bayes = recall_score(new_Y_test, preds)
      f1_bayes = f1_score(new_Y_test, preds)

      print("Accuracy:", acc_bayes)
      print("Precision:", prec_bayes)
      print("Recall:", rec_bayes)
      print("F1 Score:", f1_bayes)
```

Accuracy: 0.6976516634050881
Precision: 0.15406162464985995
Recall: 0.8870967741935484
F1 Score: 0.26252983293556087

```
[120]: draw_confusion_matrix(new_Y_test, preds, ['No Stroke', 'Stroke'])
```



CS M148 – Project 3 Report

1. Executive Summary

The main objective of this project is to predict the likelihood of a patient getting a stroke by employing the features present in the dataset. We build machine learning classification models and training them on the existing dataset.

The project is broadly divided into 4 components –

- Using descriptive statistics and basic visualizations to conduct a preliminary analysis of the data.
- Creating a pipeline strategy that allows imputation, augmentation, scaling and feature extraction from the dataset which can then be utilized for data modelling.
- Splitting and balancing the transformed data. Reducing dimensionality of the training data and finally building and optimizing the machine learning classification models using the transformed data. We output the following metrics for each model – accuracy, precision, recall, and F1 score.
- Cross validating the results obtained from the training data.

First, we display the mean, median and standard deviation of each numerical column feature in the dataset and obtain the correlation between different features to determine which of them should be retained for training the classification models.

Next, we determined whether the existing input parameters were numerical or categorical in nature, and in case of the latter decided whether it was ordinal or non-ordinal relationship. We then executed a pipeline that did the following –

- Augmented the dataset using a feature crossing and replaced null values with the respective column mean.
- Normalized the numerical features by scaling them in accordance with a standard normal distribution. Performed a one-hot encoding or label-encoding on the categorical features depending on whether it was ordinal or non-ordinal.

After, we split the transformed data into training and test sets and balance the data by generating synthetic samples of the minority class (occurrences of a stroke in this case). Then, we apply a Logistic Regression model and determine which features to prioritize by displaying some basic inferential statistics.

Next, we reduce the dimensionality of the data using PCA (Principal Components Analysis). We split and balance the data once again and then model the training data using the following classification strategies –

- Leveraging bagging on a decision tree classifier.
- Multi-Layer Perceptron classifier

We then optimize the following models using standard hyperparameter tuning techniques and then perform a cross-validation on our training results for the aforementioned models.

Finally, we implement 2 classification models of our own choice – an SVM (Support Vector Machine) and a Naïve Bayes Classifier.

2. Background/Introduction

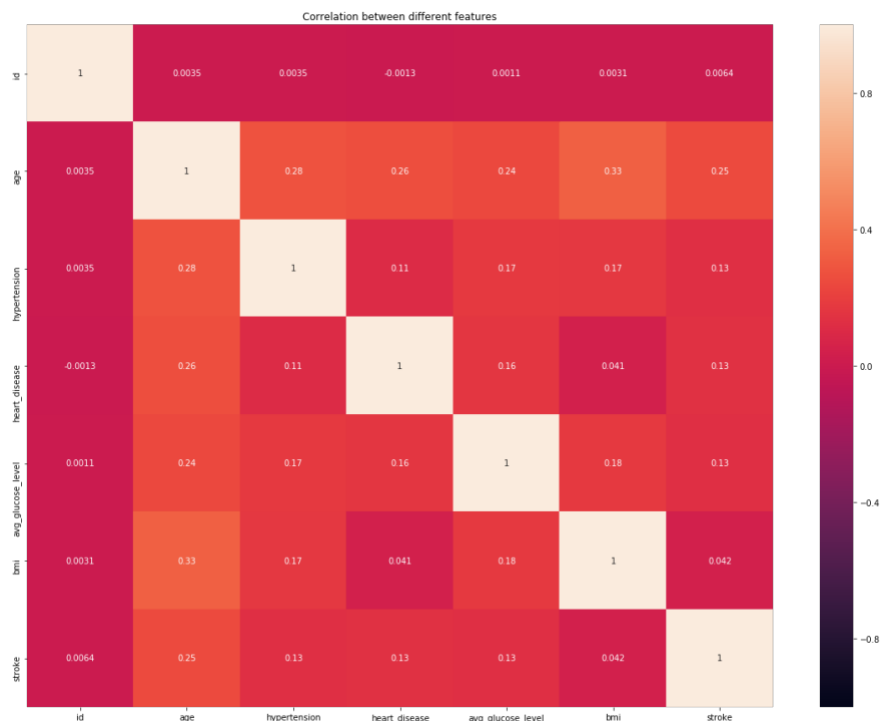
Some of the major domain challenges faced are –

- Eliminating the number of false negatives as this would reduce the usability of our model, since we would not be able to predict strokes in patients accurately.
- Assessing what features are best to model the risk of a patient getting a stroke. Some features may not easily be available as a consequence of patients being reluctant to disclose them. Hence, doctors and nurses may interpret statistics differently than classification models and thus there would be some discrepancy there.

While doctors would traditionally diagnose patients for risk of stroke themselves after looking at patient data, the introduction of machine learning in healthcare could provide new accuracy and efficiency in diagnosing patients for being at risk of strokes and thus work is needed to optimize these models.

3. Methodology

First, we do some basic statistical analysis - obtain the mean, median and standard deviation of each numerical column feature in the dataset and obtain the correlation between different features using a heatmap (shown below) to determine which of them should be retained for training the classification models.



Since `ever_married` and `Residence_type` are ordinal and binary categorical variables we perform a label-encoding on them prior to implementing the pipeline.

For the pipeline –

- We scale the three numerical features – bmi, average glucose level, and age – using StandardScaler() for better model performance and to ensure when we perform PCA later it does not skew towards high magnitude features.
- We augment a new feature **bmi_x_age** by multiplying the columns of bmi and age in order since those have the highest correlation amongst any two features and were positively correlated w.r.t stroke.
- Since only bmi has null values, we simply impute it with the column mean as it is the most representative of its distribution.
- The 3 non-ordinal categorical variables - smoking_status, work_type, and gender are one-hot encoded so that the impact for each unique value for each label is captured in our models.

We split the dataset into test and training sets after the pipeline transformation in order to model it.

It is also necessary to balance the data because of the skewed nature of the target label – 4861 labels correspond to having no stroke and 249 correspond to having a stroke – implying the dataset is highly imbalanced. Since using SMOTE to generate synthetic labels avoids overfitting as seen in the case with oversampling, we choose that particular method to balance it. We only balance the training data to make sure the testing set is representative of the original imbalanced data. If not, it can lead to our models being skewed and inefficient when dealing with real-world data.

Our first model is a logistic regression – we use both ‘sag’ and ‘lbfgs’ solvers to see if there is any difference in the two.

We then perform a PCA on the transformed data obtained via the pipeline, setting the number of components parameter to 0.8 in order to obtain the minimum number of components so that 80% of the variance is retained. This is to make sure there is no overfitting. We then split and balance the data in the same manner as described previously. We perform logistic regression again and observe the difference in results.

For optimizing our results in the following 2 models, we focus on F1 score since the target class is highly imbalanced.

- For our ensemble method, we leverage bagging with a decision tree classifier and optimize it using hyperparameter tuning on the following inputs – n_estimators, max_samples, and max_features.
- For our Neural Network classifier, we use a MLP classifier and optimize it using hyperparameter tuning on the following inputs - hidden_layer_sizes, activation, solver, alpha, and learning_rate.

We then cross-validate our training results based on F1 score using 2 approaches – StratifiedShuffleSplit and StratifiedKFold – both of which ensure that the training and test sets are representative of the actual data.

Finally, we implement two of our models – SVM and a Naïve Bayes Classifier.

4. Results

Logistic Regression: (before PCA transformation)

‘sag’:

Accuracy: 0.7534246575342466
Precision: 0.1701388888888889
Recall: 0.7903225806451613
F1 Score: 0.27999999999999997

‘lbfgs’:

Accuracy: 0.7524461839530333
Precision: 0.1695501730103806
Recall: 0.7903225806451613
F1 Score: 0.2792022792022792

Logistic Regression after PCA transformation:

‘sag’:

Accuracy: 0.7397260273972602
Precision: 0.17096774193548386
Recall: 0.8548387096774194
F1 Score: 0.2849462365591398

‘lbfgs’:

Accuracy: 0.7397260273972602
Precision: 0.17096774193548386
Recall: 0.8548387096774194
F1 Score: 0.2849462365591398

Bagging(Using a decision tree classifier):

Non-PCA data:

Accuracy: 0.7397260273972602
Precision: 0.16883116883116883
Recall: 0.8387096774193549
F1 Score: 0.28108108108108104

PCA data:

Accuracy: 0.700587084148728
Precision: 0.1534090909090909
Recall: 0.8709677419354839
F1 Score: 0.2608695652173913

Hyperparameter tuning for F1 score:

Best: 0.936358 using {'max_features': 4, 'max_samples': 1.0, 'n_estimators': 100}

MLP classifier:

PCA data:

Accuracy: 0.8424657534246576

Precision: 0.14388489208633093

Recall: 0.3225806451612903

F1 Score: 0.19900497512437812

Hyperparameter tuning for F1 score:

Best: 0.924051 using {'activation': 'relu', 'alpha': 0.0001, 'hidden_layer_sizes': (50, 100, 150), 'learning_rate': 'adaptive', 'solver': 'adam'}

KFold cross-validation:

StratifiedKFold:

Bagging f1 score: 50.55%

Neural Network f1 score: 53.51%

StratifiedShuffleSplit:

Bagging f1 score: 49.84%

Neural Network f1 score: 52.10%

SVM:

Accuracy: 0.735812133072407

Precision: 0.15100671140939598

Recall: 0.7258064516129032

F1 Score: 0.25

Naïve Bayes:

Accuracy: 0.6976516634050881

Precision: 0.15406162464985995

Recall: 0.8870967741935484

F1 Score: 0.26252983293556087

5. Discussion

We evaluate our models based primarily on F1 score due to the imbalanced nature of our dataset. Evaluating solely in accuracy may cause us to have skewed results towards not predicting a stroke and thus will not be useful in the medical field since we are unable to predict strokes in cases where patients actually have them.

The best model (prior to hyperparameter tuning) using this criterion is Logistic Regression using PCA-transformed data.

The best model (after hyperparameter tuning) using this criterion is Bagging while leveraging a Decision Tree Classifier using PCA-transformed data, with the parameters as in the results.

For all models, we see the F1 score is between 0.25-0.30 and the accuracy usually ranges between 0.75-0.85. This is because the dataset is highly imbalanced with few cases of stroke, so the model is able to obtain high accuracy by simply predicting a non-stroke for a lot of test cases, but the few number of labels in which stroke is incident causes the model to have a large number of false positives compared to true positives and thus a low precision and therefore a low F1 score.

Recommendations for Hospitals:

- The UCLA Hospital can use any of the existing models developed in the project by simply inputting the features they have to classify whether or not a patient has risk of suffering a stroke.
- Since the testing and training data were obtained after transforming the data using pipeline, the hospital may need to modify the pipeline according to the features they have in order to make sure the data is transformed in such a manner so that models can interpret it. This shouldn't be too complicated – they just have to determine whether the features they input are numerical or categorical (ordinal or non-ordinal) and with some slight modifications the pipeline can be implemented.
- In this project, I split the transformed data before balancing it so that the testing set was representative of the original data. The data may be different in the hospital so they can reverse the two steps accordingly in order to improve the F1 score though it must be kept in mind that balancing the dataset before can lead to inefficient results on real-world data as it is not representative. If the data is balanced, then using SMOTE or any other balancing technique can be avoided.

6. Conclusion

In this project, we have highlighted the difficulty of predicting a stroke due to the imbalanced nature of the dataset as a consequence of real-world statistics but has given us relatively efficient classification models for the same. We have developed models that have high accuracy as well as high recall but low F1 score— these can be used to determine if a patient is at no risk for suffering a stroke as well as predict the incidence of stroke quite well but eschews a high number of false positives. We are able to lower the dimensionality of the training set in order to prevent the risk of overfitting due to a high number of features and eventually choose the

best-suited model. We finally performed a KFold cross validation which determined that our results were generalizable to real-world data.