

Project1-Main

April 13, 2021

0.1 Introduction

Welcome to **CS188 - Data Science Fundamentals!** This course is designed to equip you with the tools and experiences necessary to start you off on a life-long exploration of datascience. We do not assume a prerequisite knowledge or experience in order to take the course.

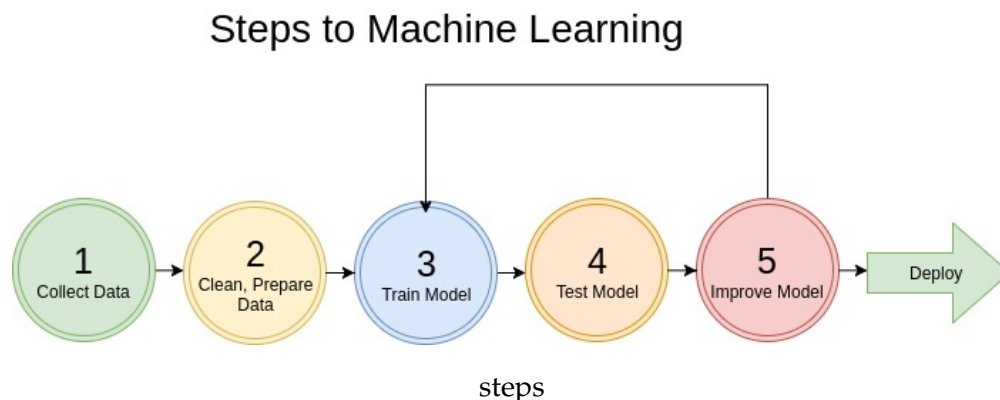
For this first project we will introduce you to the end-to-end process of doing a datascience project. Our goals for this project are to:

1. Familiarize you with the development environment for doing datascience
2. Get you comfortable with the python coding required to do datascience
3. Provide you with an sample end-to-end project to help you visualize the steps needed to complete a project on your own
4. Ask you to recreate a similar project on a separate dataset

In this project you will work through an example project end to end. Many of the concepts you will encounter will be unclear to you. That is OK! The course is designed to teach you these concepts in further detail. For now our focus is simply on having you replicate the code successfully and seeing a project through from start to finish.

Here are the main steps:

1. Get the data
2. Visualize the data for insights
3. Preprocess the data for your machine learning algorithm
4. Select a model and train
5. Does it meet the requirements? Fine tune the model



0.2 Working with Real Data

It is best to experiment with real-data as opposed to artificial datasets.

There are many different open datasets depending on the type of problems you might be interested in!

Here are a few data repositories you could check out: - [UCI Datasets](#) - [Kaggle Datasets](#) - [AWS Datasets](#)

0.3 Submission Instructions

When you have completed this assignment please save the notebook as a PDF file and submit the assignment via Gradescope

1 Example Datascience Exercise

Below we will run through an California Housing example collected from the 1990's.

1.1 Setup

```
[1]: import sys
assert sys.version_info >= (3, 5) # python>=3.5
import sklearn
assert sklearn.__version__ >= "0.20" # sklearn >= 0.20

import numpy as np #numerical package in python
import os
%matplotlib inline
import matplotlib.pyplot as plt #plotting package

# to make this notebook's output identical at every run
np.random.seed(42)

#matplotlib magic for inline figures
%matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

# Where to save the figures
ROOT_DIR = "."
IMAGES_PATH = os.path.join(ROOT_DIR, "images")
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_name, tight_layout=True, fig_extension="png", resolution=300):
    """
    plt.savefig wrapper. refer to
    https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.savefig.html
```

```

    Args:
        fig_name (str): name of the figure
        tight_layout (bool): adjust subplot to fit in the figure area
        fig_extension (str): file format to save the figure in
        resolution (int): figure resolution
    """
    path = os.path.join(IMAGES_PATH, fig_name + "." + fig_extension)
    print("Saving figure", fig_name)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

```

```

[2]: import os
import tarfile
import urllib
DATASET_PATH = os.path.join("datasets", "housing")

```

1.2 Step 1. Getting the data

1.2.1 Intro to Data Exploration Using Pandas

In this section we will load the dataset, and visualize different features using different types of plots.

Packages we will use: - **Pandas**: is a fast, flexible and expressive data structure widely used for tabular and multidimensional datasets. - **Matplotlib**: is a 2d python plotting library which you can use to create quality figures (you can plot almost anything if you're willing to code it out!) - other plotting libraries: [seaborn](#), [ggplot2](#)

```

[3]: import pandas as pd

def load_housing_data(housing_path):
    """
        loads housing.csv dataset stored

    Args:
        housing_path (str): path to folder containing housing dataset

    Returns:
        pd.DataFrame
    """
    csv_path = os.path.join(housing_path, "housing.csv")
    return pd.read_csv(csv_path)

```

```

[4]: pd.DataFrame

```

```

[4]: pandas.core.frame.DataFrame

```

```

[5]: housing = load_housing_data(DATASET_PATH) # we load the pandas dataframe
housing.head() # show the first few elements of the dataframe
              # typically this is the first thing you do

```

```
# to see how the dataframe looks like
```

```
[5]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 880.0 129.0
1 -122.22 37.86 21.0 7099.0 1106.0
2 -122.24 37.85 52.0 1467.0 190.0
3 -122.25 37.85 52.0 1274.0 235.0
4 -122.25 37.85 52.0 1627.0 280.0

population households median_income median_house_value ocean_proximity
0 322.0 126.0 8.3252 452600.0 NEAR BAY
1 2401.0 1138.0 8.3014 358500.0 NEAR BAY
2 496.0 177.0 7.2574 352100.0 NEAR BAY
3 558.0 219.0 5.6431 341300.0 NEAR BAY
4 565.0 259.0 3.8462 342200.0 NEAR BAY
```

A dataset may have different types of features - real valued - Discrete (integers) - categorical (strings)

The two categorical features are essentially the same as you can always map a categorical string/character to an integer.

In the dataset example, all our features are real valued floats, except ocean proximity which is categorical.

```
[6]: # to see a concise summary of data types, null values, and counts
# use the info() method on the dataframe
housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude          20640 non-null float64
latitude           20640 non-null float64
housing_median_age 20640 non-null float64
total_rooms         20640 non-null float64
total_bedrooms      20433 non-null float64
population          20640 non-null float64
households          20640 non-null float64
median_income       20640 non-null float64
median_house_value  20640 non-null float64
ocean_proximity     20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[7]: # you can access individual columns similarly
# to accessing elements in a python dict
housing["ocean_proximity"].head() # added head() to avoid printing many columns.
→ .
```

```
[7]: 0    NEAR BAY
      1    NEAR BAY
      2    NEAR BAY
      3    NEAR BAY
      4    NEAR BAY
      Name: ocean_proximity, dtype: object
```

```
[8]: # to access a particular row we can use iloc
      housing.iloc[1]
```

```
[8]: longitude          -122.22
      latitude           37.86
      housing_median_age      21
      total_rooms           7099
      total_bedrooms         1106
      population            2401
      households            1138
      median_income          8.3014
      median_house_value     358500
      ocean_proximity        NEAR BAY
      Name: 1, dtype: object
```

```
[9]: # one other function that might be useful is
      # value_counts(), which counts the number of occurrences
      # for categorical features
      housing["ocean_proximity"].value_counts()
```

```
[9]: <1H OCEAN      9136
      INLAND      6551
      NEAR OCEAN   2658
      NEAR BAY     2290
      ISLAND        5
      Name: ocean_proximity, dtype: int64
```

```
[10]: # The describe function compiles your typical statistics for each
       # column
       housing.describe()
```

```
[10]:
```

	longitude	latitude	housing_median_age	total_rooms	\
count	20640.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	
std	2.003532	2.135952	12.585558	2181.615252	
min	-124.350000	32.540000	1.000000	2.000000	
25%	-121.800000	33.930000	18.000000	1447.750000	
50%	-118.490000	34.260000	29.000000	2127.000000	
75%	-118.010000	37.710000	37.000000	3148.000000	
max	-114.310000	41.950000	52.000000	39320.000000	

	total_bedrooms	population	households	median_income	\
count	20433.000000	20640.000000	20640.000000	20640.000000	

mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

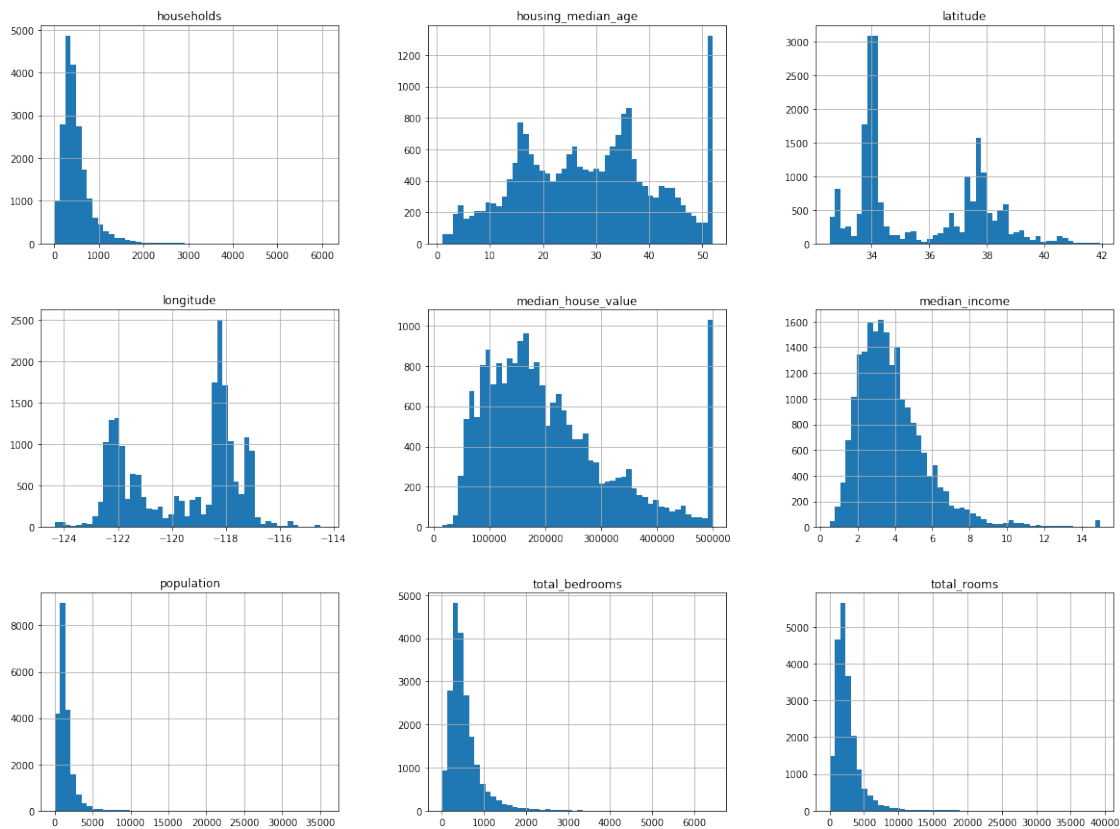
	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

If you want to learn about different ways of accessing elements or other functions it's useful to check out the getting started section [here](#)


1.3 Step 2. Visualizing the data

1.3.1 Let's start visualizing the dataset

```
[11]: # We can draw a histogram for each of the dataframes features
      # using the hist function
housing.hist(bins=50, figsize=(20,15))
      # save_fig("attribute_histogram_plots")
plt.show() # pandas internally uses matplotlib, and to display all the figures
           # the show() function must be called
```



```
[12]: # if you want to have a histogram on an individual feature:
housing["median_income"].hist()
plt.show()
```



We can convert a floating point feature to a categorical feature by binning or by defining a set of intervals.

For example, to bin the households based on median_income we can use the pd.cut function


```
[13]: # assign each bin a categorical value [1, 2, 3, 4, 5] in this case.
housing["income_cat"] = pd.cut(housing["median_income"],
                               bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
                               labels=[1, 2, 3, 4, 5])

housing["income_cat"].value_counts()
```

```
[13]: 3    7236
      2    6581
      4    3639
      5    2362
      1     822
      Name: income_cat, dtype: int64
```

```
[14]: housing["income_cat"].hist()
```

```
[14]: <matplotlib.axes._subplots.AxesSubplot at 0xa11751f98>
```




Next let's visualize the household incomes based on latitude & longitude coordinates

```
[15]: ## here's a not so interesting way of plotting it
housing.plot(kind="scatter", x="longitude", y="latitude")
```




```
save_fig("bad_visualization_plot")
```

Saving figure bad_visualization_plot



```
[16]: # we can make it look a bit nicer by using the alpha parameter,  
# it simply plots less dense areas lighter.  
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)  
save_fig("better_visualization_plot")
```

Saving figure better_visualization_plot



```
[17]: # A more interesting plot is to color code (heatmap) the dots
# based on income. The code below achieves this

# load an image of california
images_path = os.path.join('./', "images")
os.makedirs(images_path, exist_ok=True)
filename = "california.png"

import matplotlib.image as mpimg
california_img=mpimg.imread(os.path.join(images_path, filename))
ax = housing.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),
                  s=housing['population']/100, label="Population",
                  c="median_house_value", cmap=plt.get_cmap("jet"),
                  colorbar=False, alpha=0.4,
                  )

# overlay the california map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(california_img, extent=[-124.55, -113.80, 32.45, 42.05], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

# setting up heatmap colors based on median_house_value feature
```

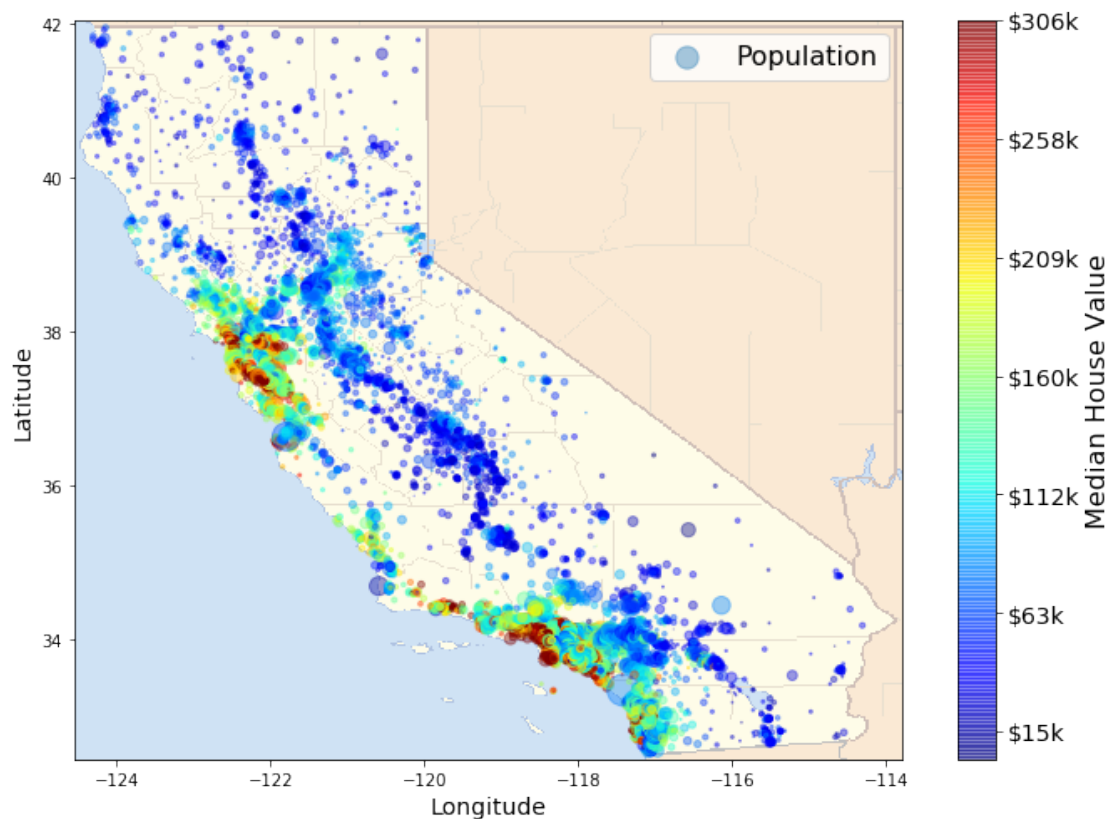
```

prices = housing["median_house_value"]
tick_values = np.linspace(prices.min(), prices.max(), 11)
cb = plt.colorbar()
cb.ax.set_yticklabels(["$%dk"%(round(v/1000)) for v in tick_values],
    ↳fontsize=14)
cb.set_label('Median House Value', fontsize=16)

plt.legend(fontsize=16)
save_fig("california_housing_prices_plot")
plt.show()

```

Saving figure california_housing_prices_plot



Not surprisingly, we can see that the most expensive houses are concentrated around the San Francisco/Los Angeles areas.

Up until now we have only visualized feature histograms and basic statistics.

When developing machine learning models the predictiveness of a feature for a particular target of interest is what's important.

It may be that only a few features are useful for the target at hand, or features may need to be augmented by applying certain transformations.

None the less we can explore this using correlation matrices. If you need to brush up on correlation take a look [here](#).

```

[18]: corr_matrix = housing.corr() # compute the correlation matrix

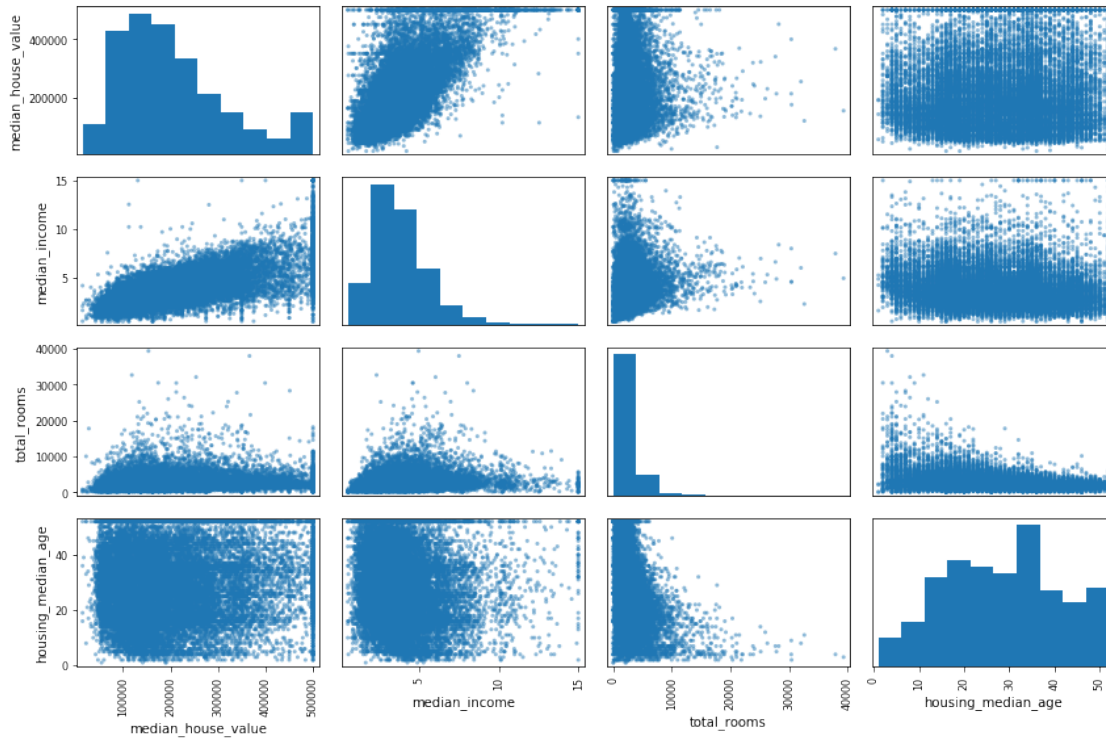
[19]: # for example if the target is "median_house_value", most correlated features
      ↪ can be sorted
      # which happens to be "median_income". This also intuitively makes sense.
      corr_matrix["median_house_value"].sort_values(ascending=False)

[19]: median_house_value    1.000000
      median_income        0.688075
      total_rooms          0.134153
      housing_median_age    0.105623
      households           0.065843
      total_bedrooms        0.049686
      population           -0.024650
      longitude            -0.045967
      latitude             -0.144160
      Name: median_house_value, dtype: float64

[20]: # the correlation matrix for different attributes/features can also be plotted
      # some features may show a positive correlation/negative correlation or
      # it may turn out to be completely random!
      from pandas.plotting import scatter_matrix
      attributes = ["median_house_value", "median_income", "total_rooms",
                   "housing_median_age"]
      scatter_matrix(housing[attributes], figsize=(12, 8))
      save_fig("scatter_matrix_plot")


```

Saving figure scatter_matrix_plot



```
[21]: # median income vs median house value plot plot 2 in the first row of top figure
housing.plot(kind="scatter", x="median_income", y="median_house_value",
              alpha=0.1)
plt.axis([0, 16, 0, 550000])
save_fig("income_vs_house_value_scatterplot")
```

Saving figure income_vs_house_value_scatterplot



1.3.2 Augmenting Features

New features can be created by combining different columns from our data set.

- $\text{rooms_per_household} = \text{total_rooms} / \text{households}$
- $\text{bedrooms_per_room} = \text{total_bedrooms} / \text{total_rooms}$
- etc.


```
[22]: housing["rooms_per_household"] = housing["total_rooms"]/housing["households"]
housing["bedrooms_per_room"] = housing["total_bedrooms"]/housing["total_rooms"]
housing["population_per_household"] = housing["population"]/housing["households"]
```

```
[23]: # obtain new correlations
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```


```
[23]: median_house_value    1.000000
median_income             0.688075
rooms_per_household       0.151948
total_rooms               0.134153
housing_median_age        0.105623
households                0.065843
total_bedrooms            0.049686
population_per_household  -0.023737
population                -0.024650
```

```
longitude          -0.045967
latitude           -0.144160
bedrooms_per_room  -0.255880
Name: median_house_value, dtype: float64
```

```
[24]: housing.plot(kind="scatter", x="rooms_per_household", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



```
[25]: housing.plot(kind="scatter", x="median_income", y="median_house_value",
                    alpha=0.2)
plt.axis([0, 5, 0, 520000])
plt.show()
```



```
[26]: housing.describe()
```

```
[26]:
```

	longitude	latitude	housing_median_age	total_rooms \
count	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081
std	2.003532	2.135952	12.585558	2181.615252
min	-124.350000	32.540000	1.000000	2.000000
25%	-121.800000	33.930000	18.000000	1447.750000
50%	-118.490000	34.260000	29.000000	2127.000000
75%	-118.010000	37.710000	37.000000	3148.000000
max	-114.310000	41.950000	52.000000	39320.000000

	total_bedrooms	population	households	median_income \
count	20433.000000	20640.000000	20640.000000	20640.000000
mean	537.870553	1425.476744	499.539680	3.870671
std	421.385070	1132.462122	382.329753	1.899822
min	1.000000	3.000000	1.000000	0.499900
25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value	rooms_per_household	bedrooms_per_room \
count	20640.000000	20640.000000	20433.000000
mean	206855.816909	5.429000	0.213039
std	115395.615874	2.474173	0.057983

min	14999.000000	0.846154	0.100000
25%	119600.000000	4.440716	0.175427
50%	179700.000000	5.229129	0.203162
75%	264725.000000	6.052381	0.239821
max	500001.000000	141.909091	1.000000

	population_per_household
count	20640.000000
mean	3.070655
std	10.386050
min	0.692308
25%	2.429741
50%	2.818116
75%	3.282261
max	1243.333333

1.4 Step 3. Preprocess the data for your machine learning algorithm

Once we've visualized the data, and have a certain understanding of how the data looks like. It's time to clean!

Most of your time will be spent on this step, although the datasets used in this project are relatively nice and clean... in the real world it could get real dirty.

After having cleaned your dataset you're aiming for: - train set - test set

In some cases you might also have a validation set as well for tuning hyperparameters (don't worry if you're not familiar with this term yet..)

In supervised learning setting your train set and test set should contain (**feature**, **target**) tuples.
- **feature**: is the input to your model - **target**: is the ground truth label - when target is categorical the task is a classification task - when target is floating point the task is a regression task

We will make use of [scikit-learn](#) python package for preprocessing.

Scikit learn is pretty well documented and if you get confused at any point simply look up the function/object!

1.4.1 Dealing With Incomplete Data

```
[27]: # have you noticed when looking at the dataframe summary certain rows
# contained null values? we can't just leave them as nulls and expect our
# model to handle them for us so we'll have to devise a method for dealing with
# them...
sample_incomplete_rows = housing[housing.isnull().any(axis=1)].head()
sample_incomplete_rows
```

```
[27]: longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
290      -122.16    37.77                47.0         1256.0             NaN
341      -122.17    37.75                38.0          992.0             NaN
538      -122.28    37.78                29.0         5154.0             NaN
563      -122.24    37.75                45.0          891.0             NaN
696      -122.10    37.69                41.0          746.0             NaN
```

	population	households	median_income	median_house_value	\
290	570.0	218.0	4.3750	161900.0	
341	732.0	259.0	1.6196	85100.0	
538	3741.0	1273.0	2.5762	173400.0	
563	384.0	146.0	4.9489	247100.0	
696	387.0	161.0	3.9063	178400.0	

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room	\
290	NEAR BAY	3	5.761468	NaN	
341	NEAR BAY	2	3.830116	NaN	
538	NEAR BAY	2	4.048704	NaN	
563	NEAR BAY	4	6.102740	NaN	
696	NEAR BAY	3	4.633540	NaN	

	population_per_household
290	2.614679
341	2.826255
538	2.938727
563	2.630137
696	2.403727

[28]: `sample_incomplete_rows.dropna(subset=["total_bedrooms"])` *# option 1: simply*
→ drop rows that have null values

[28]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms, population, households, median_income, median_house_value, ocean_proximity, income_cat, rooms_per_household, bedrooms_per_room, population_per_household]
Index: []

[29]: `sample_incomplete_rows.drop("total_bedrooms", axis=1)` *# option 2: drop*
→ the complete feature

[29]:

	longitude	latitude	housing_median_age	total_rooms	population	\
290	-122.16	37.77	47.0	1256.0	570.0	
341	-122.17	37.75	38.0	992.0	732.0	
538	-122.28	37.78	29.0	5154.0	3741.0	
563	-122.24	37.75	45.0	891.0	384.0	
696	-122.10	37.69	41.0	746.0	387.0	

	households	median_income	median_house_value	ocean_proximity	income_cat	\
290	218.0	4.3750	161900.0	NEAR BAY	3	
341	259.0	1.6196	85100.0	NEAR BAY	2	
538	1273.0	2.5762	173400.0	NEAR BAY	2	
563	146.0	4.9489	247100.0	NEAR BAY	4	
696	161.0	3.9063	178400.0	NEAR BAY	3	

	rooms_per_household	bedrooms_per_room	population_per_household
--	---------------------	-------------------	--------------------------

290	5.761468	NaN	2.614679
341	3.830116	NaN	2.826255
538	4.048704	NaN	2.938727
563	6.102740	NaN	2.630137
696	4.633540	NaN	2.403727

```
[30]: median = housing["total_bedrooms"].median()
sample_incomplete_rows["total_bedrooms"].fillna(median, inplace=True) # option 1
→ 3: replace na values with median values
sample_incomplete_rows
```

```
[30]: longitude latitude housing_median_age total_rooms total_bedrooms \
290 -122.16 37.77 47.0 1256.0 435.0
341 -122.17 37.75 38.0 992.0 435.0
538 -122.28 37.78 29.0 5154.0 435.0
563 -122.24 37.75 45.0 891.0 435.0
696 -122.10 37.69 41.0 746.0 435.0
```

	population	households	median_income	median_house_value	\
290	570.0	218.0	4.3750	161900.0	
341	732.0	259.0	1.6196	85100.0	
538	3741.0	1273.0	2.5762	173400.0	
563	384.0	146.0	4.9489	247100.0	
696	387.0	161.0	3.9063	178400.0	

	ocean_proximity	income_cat	rooms_per_household	bedrooms_per_room	\
290	NEAR BAY	3	5.761468	NaN	
341	NEAR BAY	2	3.830116	NaN	
538	NEAR BAY	2	4.048704	NaN	
563	NEAR BAY	4	6.102740	NaN	
696	NEAR BAY	3	4.633540	NaN	

	population_per_household
290	2.614679
341	2.826255
538	2.938727
563	2.630137
696	2.403727

Could you think of another plausible imputation for this dataset? (Not graded)

1.4.2 Prepare Data

Recall we are trying to predict the median house value, our features will contain longitude, latitude, housing_median_age... and our target will be median_house_value

```
[31]:
```

```
housing_features = housing.drop("median_house_value", axis=1) # drop labels for
    ↳ training set features
    # the input to the model
    ↳ should not contain the true label
housing_labels = housing["median_house_value"].copy()
```

[32]: housing_features.head()

```
[32]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -122.23 37.88 41.0 880.0 129.0
1 -122.22 37.86 21.0 7099.0 1106.0
2 -122.24 37.85 52.0 1467.0 190.0
3 -122.25 37.85 52.0 1274.0 235.0
4 -122.25 37.85 52.0 1627.0 280.0

population households median_income ocean_proximity income_cat \
0 322.0 126.0 8.3252 NEAR BAY 5
1 2401.0 1138.0 8.3014 NEAR BAY 5
2 496.0 177.0 7.2574 NEAR BAY 5
3 558.0 219.0 5.6431 NEAR BAY 4
4 565.0 259.0 3.8462 NEAR BAY 3

rooms_per_household bedrooms_per_room population_per_household
0 6.984127 0.146591 2.555556
1 6.238137 0.155797 2.109842
2 8.288136 0.129516 2.802260
3 5.817352 0.184458 2.547945
4 6.281853 0.172096 2.181467
```

```
[33]: # This cell implements the complete pipeline for preparing the data
# using sklearn's TransformerMixins
# Earlier we mentioned different types of features: categorical, and floats.
# In the case of floats we might want to convert them to categories.
# On the other hand categories in which are not already represented as integers
    ↳ must be mapped to integers before
# feeding to the model.

# Additionally, categorical values could either be represented as one-hot
    ↳ vectors or simple as normalized/unnormalized integers.
# Here we encode them using one hot vectors.

# DO NOT WORRY IF YOU DO NOT UNDERSTAND ALL THE STEPS OF THIS PIPELINE.
    ↳ CONCEPTS LIKE NORMALIZATION,
# ONE-HOT ENCODING ETC. WILL ALL BE COVERED IN DISCUSSION

from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
```

```

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder

from sklearn.base import BaseEstimator, TransformerMixin

imputer = SimpleImputer(strategy="median") # use median imputation for missing
→values
housing_num = housing_features.drop("ocean_proximity", axis=1) # remove the
→categorical feature
# column index
rooms_idx, bedrooms_idx, population_idx, households_idx = 3, 4, 5, 6

#
class AugmentFeatures(BaseEstimator, TransformerMixin):
    '''
        implements the previous features we had defined
        housing["rooms_per_household"] = housing["total_rooms"] /
→housing["households"]
        housing["bedrooms_per_room"] = housing["total_bedrooms"] /
→housing["total_rooms"]
        housing["population_per_household"] = housing["population"] /
→housing["households"]
    '''
    def __init__(self, add_bedrooms_per_room = True):
        self.add_bedrooms_per_room = add_bedrooms_per_room

    def fit(self, X, y=None):
        return self # nothing else to do

    def transform(self, X):
        rooms_per_household = X[:, rooms_idx] / X[:, households_idx]
        population_per_household = X[:, population_idx] / X[:, households_idx]
        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_idx] / X[:, rooms_idx]
            return np.c_[X, rooms_per_household, population_per_household,
                          bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]

attr_adder = AugmentFeatures(add_bedrooms_per_room=False)
housing_extra_attribs = attr_adder.transform(housing.values) # generate new
→features

# this will be are numirical pipeline

```

```

# 1. impute, 2. augment the feature set 3. normalize using StandardScaler()
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('attrs_adder', AugmentFeatures()),
    ('std_scaler', StandardScaler()),
])

housing_num_tr = num_pipeline.fit_transform(housing_num)

numerical_features = list(housing_num)
categorical_features = ["ocean_proximity"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

housing_prepared = full_pipeline.fit_transform(housing_features)

```

1.4.3 Splitting our dataset

First we need to carve out our dataset into a training and testing cohort. To do this we'll use `train_test_split`, a very elementary tool that arbitrarily splits the data into training and testing cohorts.

```

[34]: from sklearn.model_selection import train_test_split
data_target = housing['median_house_value']
train, test, target, target_test = train_test_split(housing_prepared,
    ↪data_target, test_size=0.3, random_state=0)

```

1.4.4 Select a model and train

Once we have prepared the dataset it's time to choose a model.

As our task is to predict the `median_house_value` (a floating value), regression is well suited for this.

```

[35]: from sklearn.linear_model import LinearRegression

lin_reg = LinearRegression()
lin_reg.fit(train, target)

# let's try the full preprocessing pipeline on a few training instances
data = test
labels = target_test

print("Predictions:", lin_reg.predict(data)[:5])
print("Actual labels:", list(labels)[:5])

```

```
Predictions: [207828.06448011 281099.80175494 176021.36890539 93643.46744928
304674.47047758]
Actual labels: [136900.0, 241300.0, 200700.0, 72500.0, 460000.0]
```

```
[36]: from sklearn.metrics import mean_squared_error

preds = lin_reg.predict(test)
mse = mean_squared_error(target_test, preds)
rmse = np.sqrt(mse)
rmse
```

```
[36]: 67879.86844243006
```

2 TODO: Applying the end-end ML steps to a different dataset.

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

3 [35 pts] Visualizing Data

3.0.1 [5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data

```
[37]: DATASET_PATH = os.path.join("datasets", "airbnb")
```

```
[38]: def load_airbnb_data(data_path):
        csv_path = os.path.join(data_path, "AB_NYC_2019.csv")
        return pd.read_csv(csv_path)
```

```
[39]: df_main = load_airbnb_data(DATASET_PATH)
df_main.head(10)
```

```
[39]:
```

	id	name	host_id	\
0	2539	Clean & quiet apt home by the park	2787	
1	2595	Skylit Midtown Castle	2845	
2	3647	THE VILLAGE OF HARLEM...NEW YORK !	4632	
3	3831	Cozy Entire Floor of Brownstone	4869	
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	
5	5099	Large Cozy 1 BR Apartment In Midtown East	7322	
6	5121	BlissArtsSpace!	7356	
7	5178	Large Furnished Room Near B'way	8967	
8	5203	Cozy Clean Guest Room - Family Apt	7490	
9	5238	Cute & Cozy Lower East Side 1 bdrm	7549	

	host_name	neighbourhood_group	neighbourhood	latitude	longitude	\
0	John	Brooklyn	Kensington	40.64749	-73.97237	
1	Jennifer	Manhattan	Midtown	40.75362	-73.98377	

2	Elisabeth	Manhattan	Harlem	40.80902	-73.94190
3	LisaRoxanne	Brooklyn	Clinton Hill	40.68514	-73.95976
4	Laura	Manhattan	East Harlem	40.79851	-73.94399
5	Chris	Manhattan	Murray Hill	40.74767	-73.97500
6	Garon	Brooklyn	Bedford-Stuyvesant	40.68688	-73.95596
7	Shunichi	Manhattan	Hell's Kitchen	40.76489	-73.98493
8	MaryEllen	Manhattan	Upper West Side	40.80178	-73.96723
9	Ben	Manhattan	Chinatown	40.71344	-73.99037

	room_type	price	minimum_nights	number_of_reviews	last_review	\
0	Private room	149	1	9	2018-10-19	
1	Entire home/apt	225	1	45	2019-05-21	
2	Private room	150	3	0	NaN	
3	Entire home/apt	89	1	270	2019-07-05	
4	Entire home/apt	80	10	9	2018-11-19	
5	Entire home/apt	200	3	74	2019-06-22	
6	Private room	60	45	49	2017-10-05	
7	Private room	79	2	430	2019-06-24	
8	Private room	79	2	118	2017-07-21	
9	Entire home/apt	150	1	160	2019-06-09	

	reviews_per_month	calculated_host_listings_count	availability_365
0	0.21	6	365
1	0.38	2	355
2	NaN	1	365
3	4.64	1	194
4	0.10	1	0
5	0.59	1	129
6	0.40	1	0
7	3.47	1	220
8	0.99	1	0
9	1.33	4	188

- pull up info on the data type for each of the data fields. Will any of these be problematic feeding into your model (you may need to do a little research on this)? Discuss:

[40]: `df_main.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
id                48895 non-null int64
name              48879 non-null object
host_id           48895 non-null int64
host_name         48874 non-null object
neighbourhood_group 48895 non-null object
neighbourhood     48895 non-null object
```



```

latitude          48895 non-null float64
longitude         48895 non-null float64
room_type         48895 non-null object
price             48895 non-null int64
minimum_nights    48895 non-null int64
number_of_reviews 48895 non-null int64
last_review       38843 non-null object
reviews_per_month 38843 non-null float64
calculated_host_listings_count 48895 non-null int64
availability_365   48895 non-null int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB

```

[Response here]

- drop the following columns: name, host_id, host_name, and last_review
- display a summary of the statistics of the loaded data

```
[41]: df_main.drop(['name', 'host_id', 'host_name', 'last_review'], axis=1,
→inplace=True)
```

```
[42]: df_main.describe()
```

```
[42]:
```

	id	latitude	longitude	price	minimum_nights \
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000

	number_of_reviews	reviews_per_month	calculated_host_listings_count \
count	48895.000000	38843.000000	48895.000000
mean	23.274466	1.373221	7.143982
std	44.550582	1.680442	32.952519
min	0.000000	0.010000	1.000000
25%	1.000000	0.190000	1.000000
50%	5.000000	0.720000	1.000000
75%	24.000000	2.020000	2.000000
max	629.000000	58.500000	327.000000



	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000


50%	45.000000
75%	227.000000
max	365.000000

3.0.2 [5 pts] Boxplot 3 features of your choice

- plot boxplots for 3 features of your choice

```
[43]: columns_chosen = ['availability_365', 'price', 'number_of_reviews']
for column in columns_chosen:
    plt.figure(figsize=(25,12))
    df_main.boxplot([column])
```





- describe what you expected to see with these features and what you actually observed

[Response here]

My Response: Expected a normal distribution of the price and number of reviews but both have a high number of outliers, probably because the dataset is significantly skewed. Availability is more evenly distributed.

High variability in price with long tail values, review numbers much more compact, however availability has a wider variance.

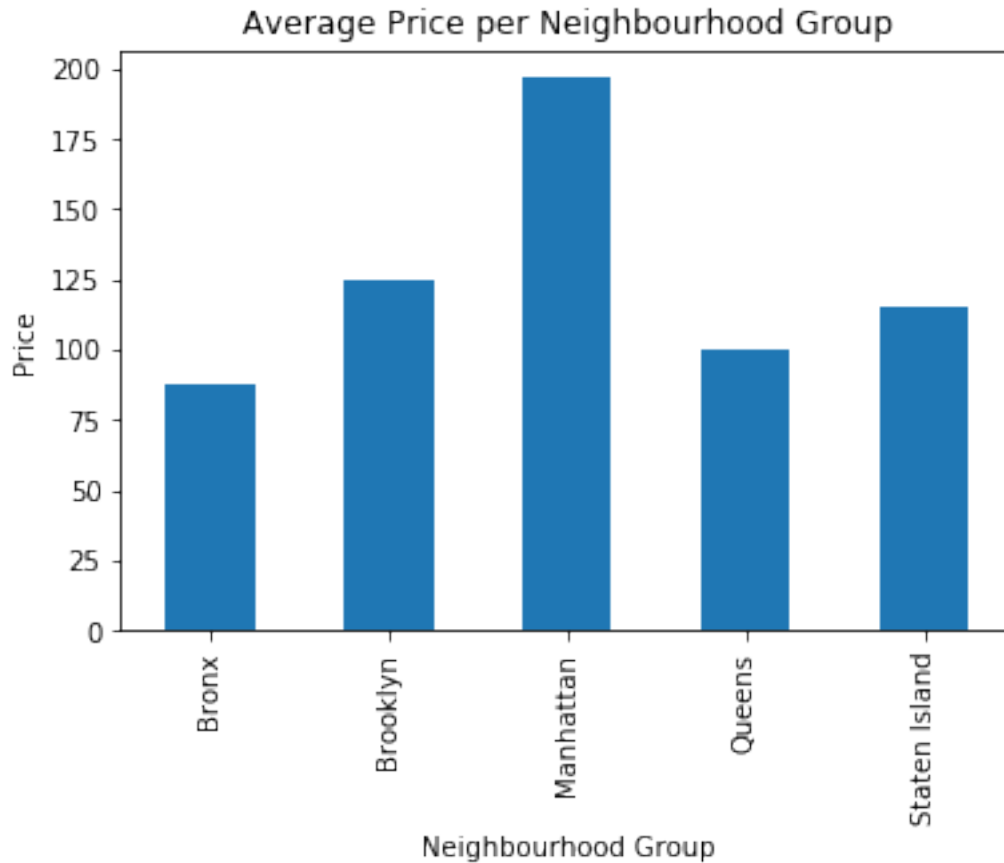
3.0.3 [10 pts] Plot average price of a listing per neighbourhood_group

```
[44]: df_main['neighbourhood_group'].value_counts()
```

```
[44]: Manhattan      21661
      Brooklyn      20104
      Queens         5666
      Bronx          1091
      Staten Island   373
      Name: neighbourhood_group, dtype: int64
```

```
[45]: df_main.groupby('neighbourhood_group')['price'].mean().plot(kind='bar')
      plt.title('Average Price per Neighbourhood Group')
      plt.xlabel('Neighbourhood Group')
      plt.ylabel('Price')
```

```
[45]: Text(0, 0.5, 'Price')
```



- describe what you expected to see with these features and what you actually observed

[Response here]

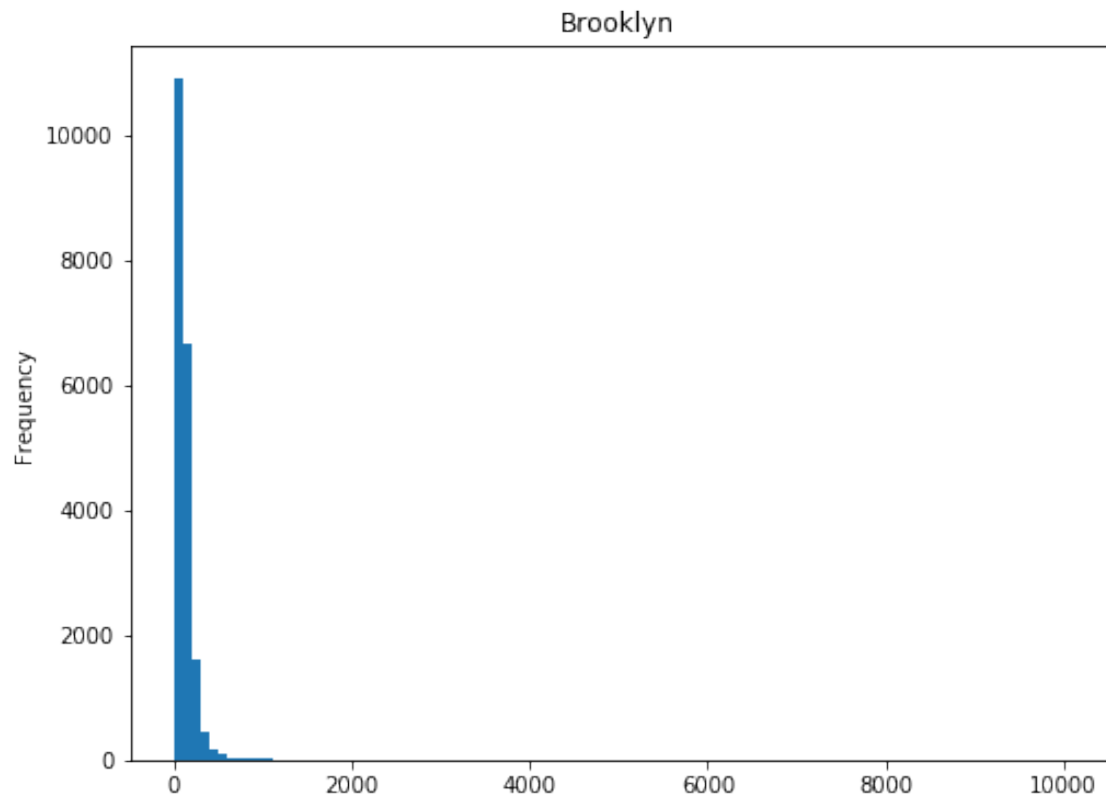
My Response: Manhattan has a higher average price than the other four, which are comparable to each other.

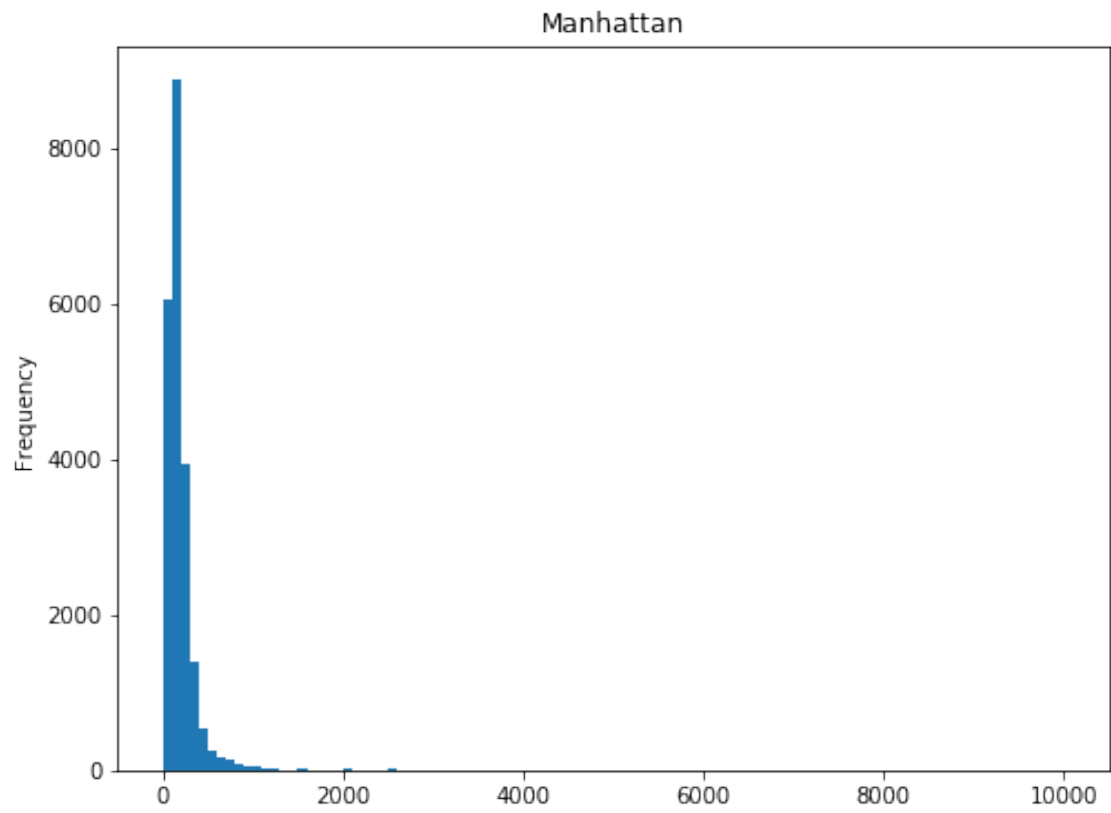
- So we can see different neighborhoods have dramatically different pricepoints, but how does the price breakdown by range. To see let's do a histogram of price by neighborhood to get a better sense of the distribution.

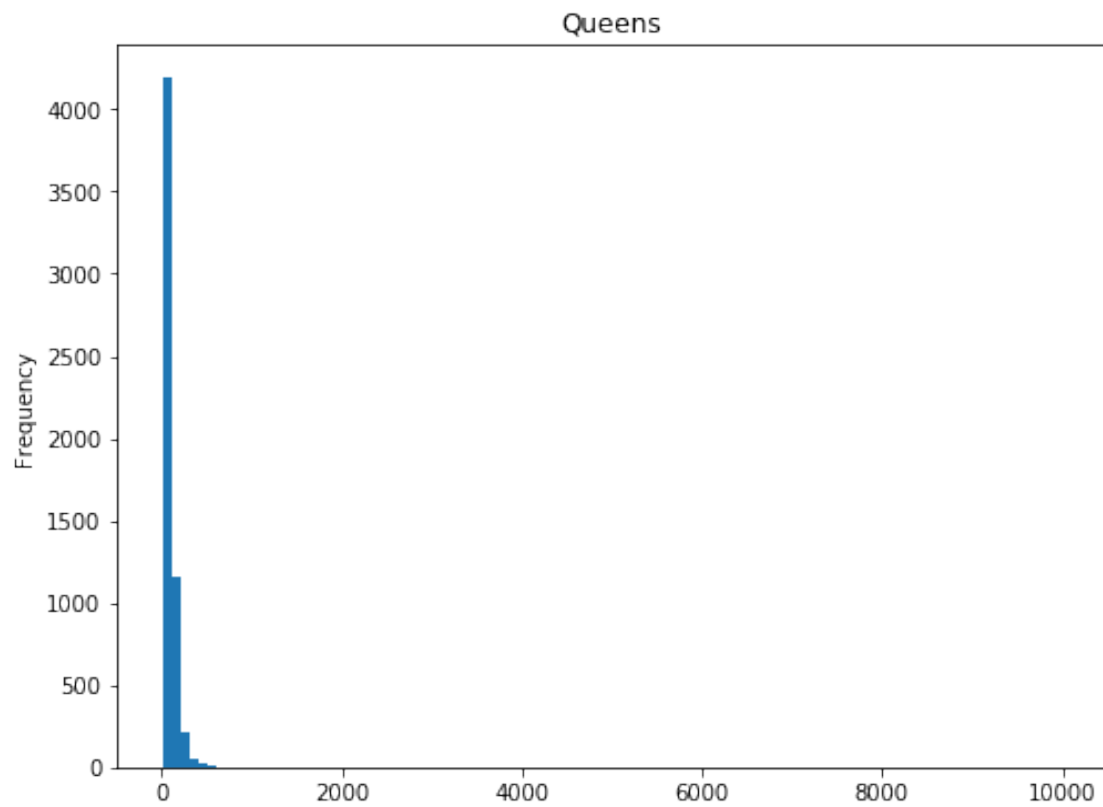
```
[46]: unique_values = df_main['neighbourhood_group'].unique()  
unique_values
```

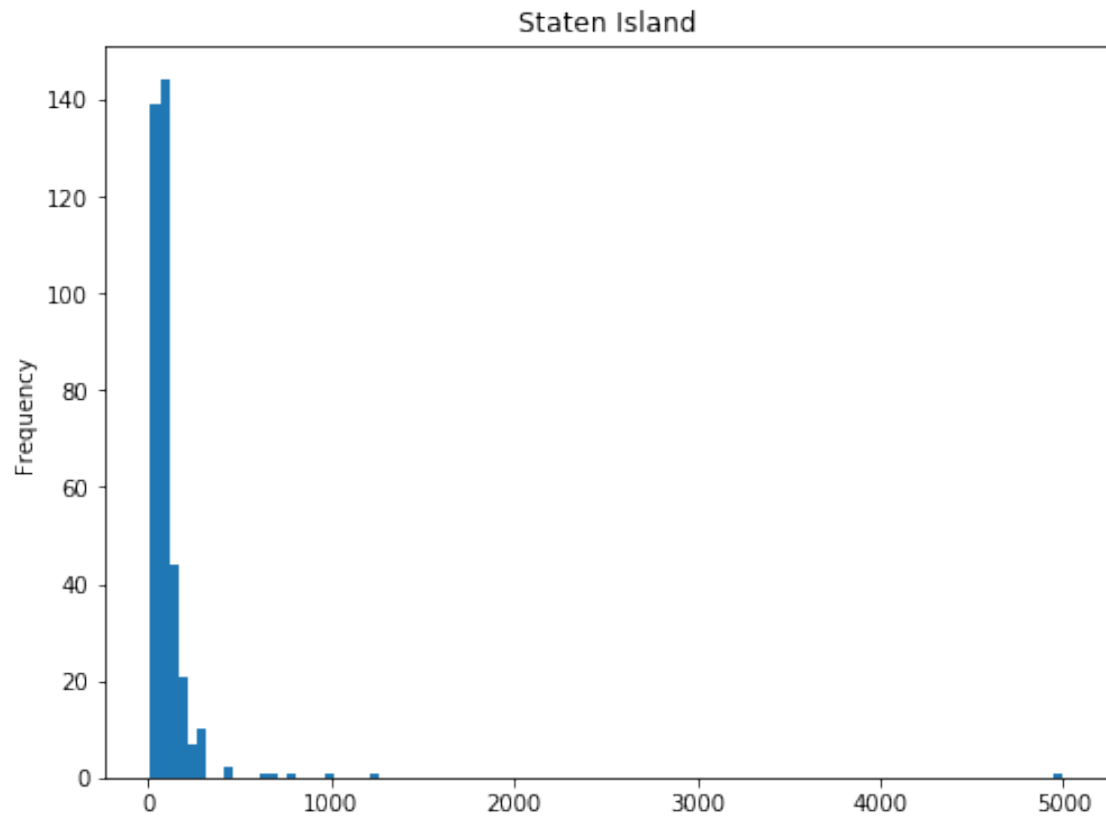
```
[46]: array(['Brooklyn', 'Manhattan', 'Queens', 'Staten Island', 'Bronx'],  
dtype=object)
```

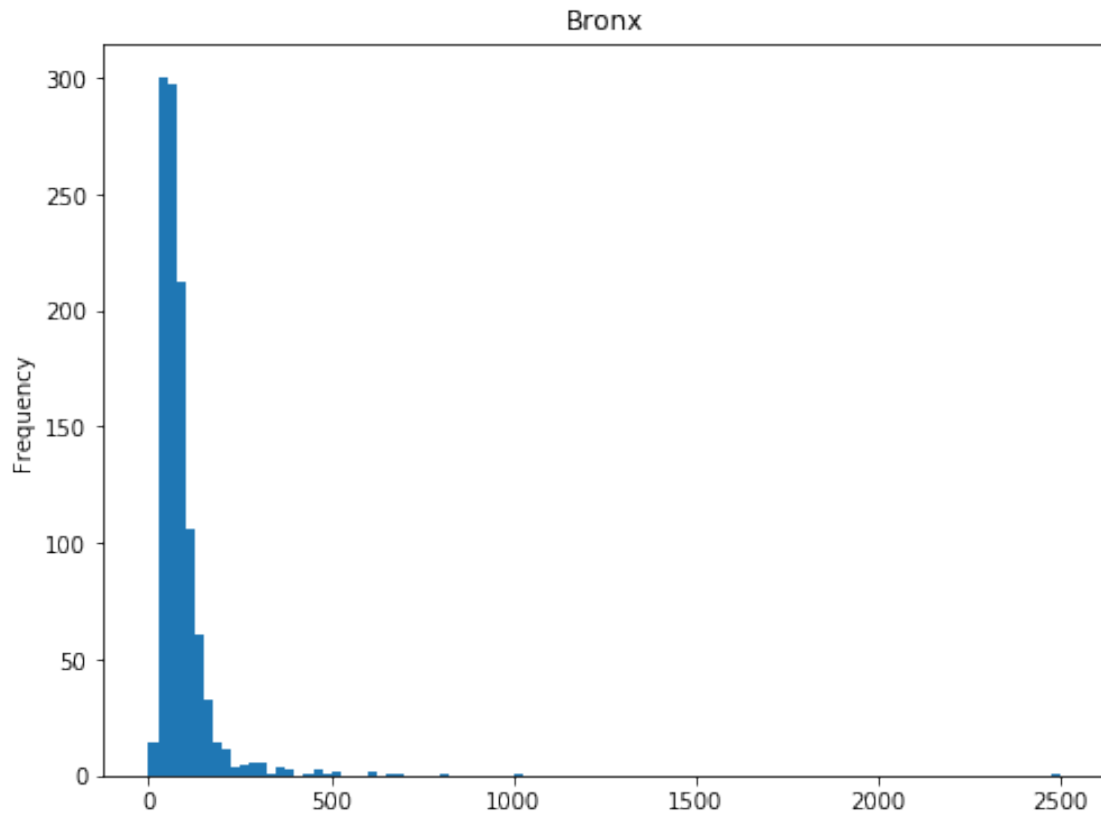
```
[47]: for n in unique_values:  
    temp_df = df_main[df_main.neighbourhood_group == n]  
    plt.figure(figsize=(8,6))  
    temp_df['price'].plot.hist(bins=100)  
    plt.title(n)
```












3.0.4 [5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```
[48]: images_path = os.path.join('./', "images")
      os.makedirs(images_path, exist_ok=True)
      filename = "newyork.png"

      import matplotlib.image as mpimg
      newyork_img=mpimg.imread(os.path.join(images_path, filename))
      ax = df_main.plot(kind="scatter", x="longitude", y="latitude", figsize=(25,12),
                          cmap=plt.get_cmap("jet"),
                          colorbar=False, alpha=0.5,
                          )
      # overlay the new york map on the plotted scatter plot

      plt.imshow(newyork_img, extent=[-74.244420, -73.712990, 40.499790, 40.913060],
                  alpha=0.5,
                  cmap=plt.get_cmap("jet"))
      plt.ylabel("Latitude", fontsize=14)
      plt.xlabel("Longitude", fontsize=14)
```

```
plt.show()
```



3.0.5 [5 pts] Plot average price of room types who have availability greater than 180 days and neighbourhood_group is Manhattan

```
[49]: temp_df = df_main[(df_main['availability_365'] > 180) &
    → (df_main['neighbourhood_group'] == 'Manhattan')]
temp_df.head()
```

```
[49]:
```

	id	neighbourhood_group	neighbourhood	latitude	longitude	\
1	2595	Manhattan	Midtown	40.75362	-73.98377	
2	3647	Manhattan	Harlem	40.80902	-73.94190	
7	5178	Manhattan	Hell's Kitchen	40.76489	-73.98493	
9	5238	Manhattan	Chinatown	40.71344	-73.99037	
13	6021	Manhattan	Upper West Side	40.79826	-73.96113	


	room_type	price	minimum_nights	number_of_reviews	\
1	Entire home/apt	225	1	45	
2	Private room	150	3	0	

7	Private room	79	2	430
9	Entire home/apt	150	1	160
13	Private room	85	2	113

	reviews_per_month	calculated_host_listings_count	availability_365
1	0.38	2	355
2	NaN	1	365
7	3.47	1	220
9	1.33	4	188
13	0.91	1	333

```
[50]: temp_df.groupby('room_type')['price'].mean().plot(kind='bar')
plt.title('Average Price per Room Type')
plt.xlabel('Room Type')
plt.ylabel('Price')
```

```
[50]: Text(0, 0.5, 'Price')
```

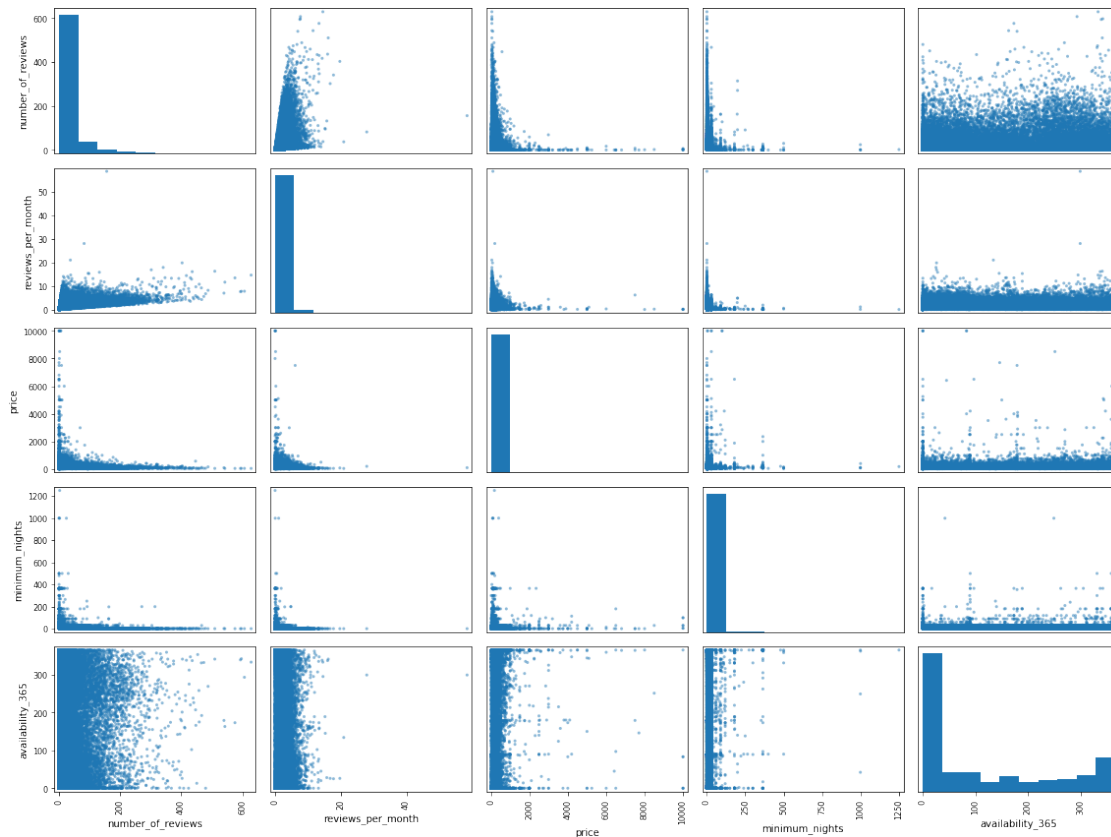


3.0.6 [5 pts] Plot correlation matrix

- which features have positive correlation?
- which features have negative correlation?

```
[51]: corr_matrix = df_main.corr()
attributes = ["number_of_reviews", "reviews_per_month", "price",
             ↪ "minimum_nights", "availability_365"]
scatter_matrix(df_main[attributes], figsize=(16, 12))
save_fig("scatter_matrix_plot2")
```

Saving figure scatter_matrix_plot2



```
[52]: corr_matrix
```

```
[52]:
```

	id	latitude	longitude	price \
id	1.000000	-0.003125	0.090908	0.010619
latitude	-0.003125	1.000000	0.084788	0.033939
longitude	0.090908	0.084788	1.000000	-0.150019
price	0.010619	0.033939	-0.150019	1.000000
minimum_nights	-0.013224	0.024869	-0.062747	0.042799
number_of_reviews	-0.319760	-0.015389	0.059094	-0.047954

reviews_per_month	0.291828	-0.010142	0.145948	-0.030608
calculated_host_listings_count	0.133272	0.019517	-0.114713	0.057472
availability_365	0.085468	-0.010983	0.082731	0.081829

	minimum_nights	number_of_reviews	\
id	-0.013224	-0.319760	
latitude	0.024869	-0.015389	
longitude	-0.062747	0.059094	
price	0.042799	-0.047954	
minimum_nights	1.000000	-0.080116	
number_of_reviews	-0.080116	1.000000	
reviews_per_month	-0.121702	0.549868	
calculated_host_listings_count	0.127960	-0.072376	
availability_365	0.144303	0.172028	

	reviews_per_month	\
id	0.291828	
latitude	-0.010142	
longitude	0.145948	
price	-0.030608	
minimum_nights	-0.121702	
number_of_reviews	0.549868	
reviews_per_month	1.000000	
calculated_host_listings_count	-0.009421	
availability_365	0.185791	

	calculated_host_listings_count	\
id	0.133272	
latitude	0.019517	
longitude	-0.114713	
price	0.057472	
minimum_nights	0.127960	
number_of_reviews	-0.072376	
reviews_per_month	-0.009421	
calculated_host_listings_count	1.000000	
availability_365	0.225701	

	availability_365
id	0.085468
latitude	-0.010983
longitude	0.082731
price	0.081829
minimum_nights	0.144303
number_of_reviews	0.172028
reviews_per_month	0.185791
calculated_host_listings_count	0.225701
availability_365	1.000000

[Response here]

In general, there does not seem to be a particularly strong correlation (positive or negative) between any pairing of these two features the exception of a (relatively) strong positive correlation reviews_per_month and number_of_reviews, which is expected. minimum_nights and price seem to have a slight positive correlation and reviews_per_month and minimum_nights seem to have a slight negative correlation but it doesn't seem apparent for any other pairings. A log transformation for some features (such as minimum_nights and price) may result a more clear correlation between pairings.

4 [30 pts] Prepare the Data

4.0.1 [5 pts] Augment the dataframe with two other features which you think would be useful

```
[53]: df_main.columns
```

```
[53]: Index(['id', 'neighbourhood_group', 'neighbourhood', 'latitude', 'longitude',  
        'room_type', 'price', 'minimum_nights', 'number_of_reviews',  
        'reviews_per_month', 'calculated_host_listings_count',  
        'availability_365'],  
        dtype='object')
```

```
[54]: #I'm assuming that the price is given per day here, couldn't find further  
      →information on it  
df_main['min_price'] = df_main['price']*df_main['minimum_nights']  
df_main['price_per_listings'] = df_main['price']/  
      →df_main['calculated_host_listings_count']
```

```
[55]: df_main.head(5)
```

```
[55]:
```

	id	neighbourhood_group	neighbourhood	latitude	longitude	\
0	2539	Brooklyn	Kensington	40.64749	-73.97237	
1	2595	Manhattan	Midtown	40.75362	-73.98377	
2	3647	Manhattan	Harlem	40.80902	-73.94190	
3	3831	Brooklyn	Clinton Hill	40.68514	-73.95976	
4	5022	Manhattan	East Harlem	40.79851	-73.94399	

	room_type	price	minimum_nights	number_of_reviews	\
0	Private room	149	1	9	
1	Entire home/apt	225	1	45	
2	Private room	150	3	0	
3	Entire home/apt	89	1	270	
4	Entire home/apt	80	10	9	

	reviews_per_month	calculated_host_listings_count	availability_365	\
0	0.21	6	365	
1	0.38	2	355	
2	NaN	1	365	
3	4.64	1	194	
4	0.10	1	0	

	min_price	price_per_listings
0	149	24.833333
1	225	112.500000
2	450	150.000000
3	89	89.000000
4	800	80.000000

4.0.2 [5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
[56]: df_main.isnull().sum()
```

```
[56]: id                                0
      neighbourhood_group              0
      neighbourhood                   0
      latitude                        0
      longitude                       0
      room_type                       0
      price                           0
      minimum_nights                  0
      number_of_reviews                0
      reviews_per_month              10052
      calculated_host_listings_count  0
      availability_365                0
      min_price                       0
      price_per_listings              0
      dtype: int64
```

```
[57]: df_main['reviews_per_month'].mean()
```

```
[57]: 1.3732214298586884
```

```
[58]: df_main['reviews_per_month'].std()
```

```
[58]: 1.6804419952744627
```

```
[62]: df_main_final = df_main.copy()
```

```
[63]: df_main_final['reviews_per_month'].fillna(df_main['reviews_per_month'].
      ↪median(), inplace=True)
```

```
[64]: df_main.head()
```

```
[64]:   id neighbourhood_group neighbourhood  latitude  longitude \
0  2539             Brooklyn   Kensington  40.64749  -73.97237
1  2595             Manhattan     Midtown  40.75362  -73.98377
2  3647             Manhattan     Harlem  40.80902  -73.94190
3  3831             Brooklyn  Clinton Hill  40.68514  -73.95976
4  5022             Manhattan   East Harlem  40.79851  -73.94399
```

	room_type	price	minimum_nights	number_of_reviews	\
0	Private room	149	1	9	
1	Entire home/apt	225	1	45	
2	Private room	150	3	0	
3	Entire home/apt	89	1	270	
4	Entire home/apt	80	10	9	

	reviews_per_month	calculated_host_listings_count	availability_365	\
0	0.21	6	365	
1	0.38	2	355	
2	NaN	1	365	
3	4.64	1	194	
4	0.10	1	0	

	min_price	price_per_listings
0	149	24.833333
1	225	112.500000
2	450	150.000000
3	89	89.000000
4	800	80.000000

```
[65]: df_main_final.head()
```

```
[65]:
```

	id	neighbourhood_group	neighbourhood	latitude	longitude	\
0	2539	Brooklyn	Kensington	40.64749	-73.97237	
1	2595	Manhattan	Midtown	40.75362	-73.98377	
2	3647	Manhattan	Harlem	40.80902	-73.94190	
3	3831	Brooklyn	Clinton Hill	40.68514	-73.95976	
4	5022	Manhattan	East Harlem	40.79851	-73.94399	

	room_type	price	minimum_nights	number_of_reviews	\
0	Private room	149	1	9	
1	Entire home/apt	225	1	45	
2	Private room	150	3	0	
3	Entire home/apt	89	1	270	
4	Entire home/apt	80	10	9	

	reviews_per_month	calculated_host_listings_count	availability_365	\
0	0.21	6	365	
1	0.38	2	355	
2	0.72	1	365	
3	4.64	1	194	
4	0.10	1	0	

	min_price	price_per_listings
0	149	24.833333
1	225	112.500000
2	450	150.000000

3	89	89.000000
4	800	80.000000

We can fill all null values of reviews per month with the median. This is because the standard deviation exceeds the mean which indicates the mean is not very representative of the data. Thus, we replace it with the median; which should not affect the overall composition of the data.

4.0.3 [15 pts] Code complete data pipeline using sklearn mixins

```
[66]: df_main_final.columns
```

```
[66]: Index(['id', 'neighbourhood_group', 'neighbourhood', 'latitude', 'longitude',
        'room_type', 'price', 'minimum_nights', 'number_of_reviews',
        'reviews_per_month', 'calculated_host_listings_count',
        'availability_365', 'min_price', 'price_per_listings'],
        dtype='object')
```

```
[71]: final_data = df_main_final.drop(["neighbourhood", "neighbourhood_group",
        ↪ "room_type", "price"], axis=1)

numerical_features = list(final_data)
categorical_features = ["neighbourhood", "neighbourhood_group", "room_type"]

full_pipeline = ColumnTransformer([
    ("num", num_pipeline, numerical_features),
    ("cat", OneHotEncoder(), categorical_features),
])

final_data_X = df_main_final.drop(columns=["price"])
final_data_Y = df_main_final["price"]
data_prepared = full_pipeline.fit_transform(final_data_X)
data_prepared
```

```
[71]: <48895x242 sparse matrix of type '<class 'numpy.float64'>'
        with 782320 stored elements in Compressed Sparse Row format>
```

4.0.4 [5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
[72]: X_train, X_test, Y_train, Y_test = train_test_split(final_data_X, final_data_Y,
        ↪ test_size=0.2, random_state=42)

print("X_train shape:", X_train.shape)
print("Y_train shape:", Y_train.shape)
print("X_test shape:", X_test.shape)
print("Y_test shape:", Y_test.shape)
```

```
X_train shape: (39116, 13)
Y_train shape: (39116,)
X_test shape: (9779, 13)
Y_test shape: (9779,)
```

5 [15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
[73]: X_train_final = X_train.drop(["neighbourhood", "neighbourhood_group",  
    ↪ "room_type"], axis=1)  
X_test_final = X_test.drop(["neighbourhood", "neighbourhood_group",  
    ↪ "room_type"], axis=1)  
  
linreg = LinearRegression()  
linreg.fit(X_train_final, Y_train)  
  
preds = linreg.predict(X_train_final)  
mse1 = mean_squared_error(Y_train, preds)  
print ("Train MSE: ", mse1)  
preds = linreg.predict(X_test_final)  
mse2 = mean_squared_error(Y_test, preds)  
print ("Test MSE: ", mse2)
```

Train MSE: 11201.842819463272

Test MSE: 11096.177490017448